

# Algerian Forest Fire

- Machine Learning LifeCycle
  1. Data Ingestion
  2. EDA
  3. Preprocessing
  4. Model Building
  5. Performance Metrics
- Problem Statement
  - The dataset includes 244 instances that regroup a data of two regions of Algeria, namely the Bejaia region located in the northeast of Algeria and the Sidi Bel-abbes region located in the northwest of Algeria.
    - 122 instances for each region.
    - The period from June 2012 to September 2012.
    - The dataset includes 11 attributes and 1 output attribute (class)
    - The 244 instances have been classified into fire (138 classes) and not fire (106 classes) classes.
- Feature Description
  1. Date : (DD/MM/YYYY) Day, month ('june' to 'september'), year (2012) Weather data observations
  2. Temp : temperature noon (temperature max) in Celsius degrees: 22 to 42
  3. RH : Relative Humidity in %: 21 to 90
  4. Ws : Wind speed in km/h: 6 to 29
  5. Rain: total day in mm: 0 to 16.8 FWI Components
  6. Fine Fuel Moisture Code (FFMC) index from the FWI system: 28.6 to 92.5
  7. Duff Moisture Code (DMC) index from the FWI system: 1.1 to 65.9
  8. Drought Code (DC) index from the FWI system: 7 to 220.4
  9. Initial Spread Index (ISI) index from the FWI system: 0 to 18.5
  10. Buildup Index (BUI) index from the FWI system: 1.1 to 68
  11. Fire Weather Index (FWI) Index: 0 to 31.1
  12. Classes: two classes, namely Fire and not Fire
- Machine Learning Models Used
  1. Linear Regression
  2. Ridge Regression
  3. Lasso Regression
  4. Elastic-Net Regression

```
In [1]: ## Importing the necessary Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings("ignore")
%matplotlib inline
```

```
In [2]: ## importing the Dataset
data=pd.read_csv(r'Algerian_Forest_Fire.csv',header=1)
```

```
In [3]: ## There are two regions Algeria and Bejaia in same file it needs to be
## converted into single dataframe for the anaylsis
data=data.drop([122,123],axis=0)
```

```
In [4]: # Printing the Column Labels in the dataset
data.columns
```

```
Out[4]: Index(['day', 'month', 'year', 'Temperature', ' RH', ' Ws', 'Rain ', 'FFMC',
'DMC', 'DC', 'ISI', 'BUI', 'FWI', 'Classes '],
dtype='object')
```

```
In [5]: ## Renaming the Columns which has space in the column labels
data.rename(columns={' RH':'RH','Rain ':'Rain',' Ws':'Ws','Classes ':'Classes'},inplace=True)
```

```
In [6]: data.columns
```

```
Out[6]: Index(['day', 'month', 'year', 'Temperature', 'RH', 'Ws', 'Rain', 'FFMC',
'DMC', 'DC', 'ISI', 'BUI', 'FWI', 'Classes'],
dtype='object')
```

```
In [7]: # Top Five Records in the dataset
data.head()
```

```
Out[7]:
```

	day	month	year	Temperature	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FWI	Classes
0	01	06	2012	29	57	18	0	65.7	3.4	7.6	1.3	3.4	0.5	not fire
1	02	06	2012	29	61	13	1.3	64.4	4.1	7.6	1	3.9	0.4	not fire
2	03	06	2012	26	82	22	13.1	47.1	2.5	7.1	0.3	2.7	0.1	not fire
3	04	06	2012	25	89	13	2.5	28.6	1.3	6.9	0	1.7	0	not fire
4	05	06	2012	27	77	16	0	64.8	3	14.2	1.2	3.9	0.5	not fire

```
In [8]: # Sample five records in the dataset
data.sample(5)
```

```
Out[8]:
```

	day	month	year	Temperature	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FWI	Classes
227	12	09	2012	31	72	14	0	84.2	8.3	25.2	3.8	9.1	3.9	fire
11	12	06	2012	26	81	19	0	84	13.8	61.4	4.8	17.7	7.1	fire
58	29	07	2012	32	73	15	0	86.6	26.7	127	5.6	35	11.9	fire
244	29	09	2012	24	54	18	0.1	79.7	4.3	15.2	1.7	5.1	0.7	not fire
18	19	06	2012	31	55	16	0.1	79.9	4.5	16	2.5	5.3	1.4	not fire

```
In [9]: # Checking whether there is any duplicate records in the dataset
data.duplicated().sum()
```

```
Out[9]: 0
```

Observation: There is no records which is duplicated in the dataset

```
In [10]: ## Checking the total number of records and columns with the dataset
data.shape
```

```
Out[10]: (244, 14)
```

# Exploratory Data Analysis

```
In [11]: ## Detailed information of each features in the dataset  
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 244 entries, 0 to 245  
Data columns (total 14 columns):  
#   Column          Non-Null Count  Dtype  
---  -  
0   day             244 non-null    object  
1   month           244 non-null    object  
2   year            244 non-null    object  
3   Temperature     244 non-null    object  
4   RH              244 non-null    object  
5   Ws              244 non-null    object  
6   Rain            244 non-null    object  
7   FFMC            244 non-null    object  
8   DMC             244 non-null    object  
9   DC              244 non-null    object  
10  ISI             244 non-null    object  
11  BUI             244 non-null    object  
12  FWI             244 non-null    object  
13  Classes         243 non-null    object  
dtypes: object(14)  
memory usage: 28.6+ KB
```

## Observation:

- There are some features which should be converted to int and float type there is some data discriminatory in the dataset which need to be preprocessed

```
In [12]: data.isnull().sum()
```

```
Out[12]: day             0  
month           0  
year            0  
Temperature     0  
RH              0  
Ws              0  
Rain            0  
FFMC            0  
DMC             0  
DC              0  
ISI             0  
BUI             0  
FWI             0  
Classes         1  
dtype: int64
```

## Observation:

- There is one record which is null in the feature Classes

```
In [13]: ## Checking the datatype of FWI feature  
data['FWI'].dtype
```

```
Out[13]: dtype('O')
```

```
In [14]: ## Checking the Shape of FWI feature  
data['FWI'].shape
```

```
Out[14]: (244,)
```

## Converting the Feature Datatype

```
In [15]: ### Rain,FFMC,DMC,DC,ISI,BUI,FWI has object type  
float_var=['Rain','FFMC','DMC','DC','ISI','BUI','FWI']
```

```
In [16]: ### All the features which are in float_var are converted to float type  
for col in float_var:  
    data[col] = data[col].apply(pd.to_numeric, downcast='float', errors='coerce')
```

```
In [17]: data.dtypes
```

```
Out[17]: day           object  
month          object  
year           object  
Temperature     object  
RH             object  
Ws             object  
Rain           float32  
FFMC           float32  
DMC            float32  
DC             float32  
ISI            float32  
BUI            float32  
FWI            float32  
Classes        object  
dtype: object
```

```
In [18]: ### day,month,year,Temperature,RH,Ws are in object type  
num_var=['day','month','year','Temperature','RH','Ws']
```

```
In [19]: ### All the features in num_var are converted into int type  
for col in num_var:  
    data[col]=data[col].astype('int')
```

```
In [20]: data.dtypes
```

```
Out[20]: day                int32
month                int32
year                int32
Temperature          int32
RH                  int32
Ws                  int32
Rain                float32
FFMC                float32
DMC                float32
DC                 float32
ISI                float32
BUI                float32
FWI                float32
Classes             object
dtype: object
```

```
In [21]: data.head()
```

Out[21]:

	day	month	year	Temperature	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FWI	Classes
0	1	6	2012	29	57	18	0.0	65.699997	3.4	7.6	1.3	3.4	0.5	not fire
1	2	6	2012	29	61	13	1.3	64.400002	4.1	7.6	1.0	3.9	0.4	not fire
2	3	6	2012	26	82	22	13.1	47.099998	2.5	7.1	0.3	2.7	0.1	not fire
3	4	6	2012	25	89	13	2.5	28.600000	1.3	6.9	0.0	1.7	0.0	not fire
4	5	6	2012	27	77	16	0.0	64.800003	3.0	14.2	1.2	3.9	0.5	not fire

Statistical Analysis

```
In [22]: ### Storing All the numerical variable in num_col
num_col=data[data.dtypes[(data.dtypes=='int') | (data.dtypes=='float')|(data.dtypes=='f...
```

```
In [23]: ### Describing the stastical analysis of the data
num_col.describe().T
```

Out[23]:

	count	mean	std	min	25%	50%	75%	max
day	244.0	15.754098	8.825059	1.0	8.000000	16.000000	23.000000	31.000000
month	244.0	7.500000	1.112961	6.0	7.000000	7.500000	8.000000	9.000000
year	244.0	2012.000000	0.000000	2012.0	2012.000000	2012.000000	2012.000000	2012.000000
Temperature	244.0	32.172131	3.633843	22.0	30.000000	32.000000	35.000000	42.000000
RH	244.0	61.938525	14.884200	21.0	52.000000	63.000000	73.250000	90.000000
Ws	244.0	15.504098	2.810178	6.0	14.000000	15.000000	17.000000	29.000000
Rain	244.0	0.760656	1.999407	0.0	0.000000	0.000000	0.500000	16.799999
FFMC	244.0	77.887680	14.337568	28.6	72.075003	83.500000	88.300003	96.000000
DMC	244.0	14.673360	12.368040	0.7	5.800000	11.300000	20.750000	65.900002
DC	243.0	49.430847	47.665600	6.9	12.350000	33.099998	69.099998	220.399994
ISI	244.0	4.774180	4.175318	0.0	1.400000	3.500000	7.300000	19.000000
BUI	244.0	16.664747	14.204824	1.1	6.000000	12.250000	22.525000	68.000000
FWI	243.0	7.035391	7.440567	0.0	0.700000	4.200000	11.450001	31.100000

### Observation

1. Rain Feature is highly skewed on the right
2. year feature is not required as it has only one unique value through out the dataset
3. Rain Features has some outliers
4. BUI also right skewed and has some outliers

### Drop the year Feature

```
In [24]: ### Unique values in the year featue  
num_col['year'].unique()
```

```
Out[24]: array([2012])
```

### Observation:

- As there is only one year it has been analysed for both the regions it can be dropped from the dataset

```
In [25]: ### Dropping the year feature  
num_col.drop(labels=['year'],axis=1,inplace=True)
```

```
In [26]: num_col.head()
```

```
Out[26]:
```

	day	month	Temperature	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FWI
0	1	6	29	57	18	0.0	65.699997	3.4	7.6	1.3	3.4	0.5
1	2	6	29	61	13	1.3	64.400002	4.1	7.6	1.0	3.9	0.4
2	3	6	26	82	22	13.1	47.099998	2.5	7.1	0.3	2.7	0.1
3	4	6	25	89	13	2.5	28.600000	1.3	6.9	0.0	1.7	0.0
4	5	6	27	77	16	0.0	64.800003	3.0	14.2	1.2	3.9	0.5

### Correlation between the features

```
In [27]: num_col.corr()
```

Out[27]:

	day	month	Temperature	RH	Ws	Rain	FFMC	DMC
day	1.000000e+00	2.232788e-17	0.095772	-0.074209	0.047001	-0.112265	0.224032	0.491577
month	2.232788e-17	1.000000e+00	-0.059017	-0.037884	-0.041447	0.035322	0.015577	0.068178
Temperature	9.577222e-02	-5.901677e-02	1.000000	-0.654443	-0.278132	-0.326786	0.677491	0.483105
RH	-7.420934e-02	-3.788419e-02	-0.654443	1.000000	0.236084	0.222968	-0.645658	-0.405133
Ws	4.700086e-02	-4.144673e-02	-0.278132	0.236084	1.000000	0.170169	-0.163255	-0.001246
Rain	-1.122654e-01	3.532207e-02	-0.326786	0.222968	0.170169	1.000000	-0.544045	-0.288548
FFMC	2.240321e-01	1.557669e-02	0.677491	-0.645658	-0.163255	-0.544045	1.000000	0.602391
DMC	4.915710e-01	6.817778e-02	0.483105	-0.405133	-0.001246	-0.288548	0.602391	1.000000
DC	5.279524e-01	1.265111e-01	0.376284	-0.226941	0.079135	-0.298023	0.507397	0.875929
ISI	1.777266e-01	6.168012e-02	0.607551	-0.690637	0.015248	-0.347105	0.739730	0.674499
BUI	5.172292e-01	8.582162e-02	0.455504	-0.348587	0.029756	-0.299171	0.589652	0.982070
FWI	3.507809e-01	8.263884e-02	0.566670	-0.580957	0.032368	-0.324422	0.691132	0.875864

#### Observation:

- Temperature and Fine Fuel Moisture Code (FFMC) are Positively Correlated .
- Temperature and Relative Humidity (RH) are Negatively Correlated .
- Temperature and Initial Spread Index (ISI) are Positively Correlated .
- Buildup Index (BUI) and Duff Moisture Code (DMC) are highly Correlated but both of them are independent variable which leads to MultiCollinearity Which should be handled
- Drought Code (DC) and Buildup Index (BUI) are also highly Correlated they also make MultiCollinearity

```
In [28]: ### Make a Copy of the DataFrame
df=num_col.copy()
```

```
In [29]: ### Appending the Classes feature in df
df['Classes']=data['Classes']
```

```
In [30]: df.head()
```

```
Out[30]:
```

	day	month	Temperature	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FWI	Classes
0	1	6	29	57	18	0.0	65.699997	3.4	7.6	1.3	3.4	0.5	not fire
1	2	6	29	61	13	1.3	64.400002	4.1	7.6	1.0	3.9	0.4	not fire
2	3	6	26	82	22	13.1	47.099998	2.5	7.1	0.3	2.7	0.1	not fire
3	4	6	25	89	13	2.5	28.600000	1.3	6.9	0.0	1.7	0.0	not fire
4	5	6	27	77	16	0.0	64.800003	3.0	14.2	1.2	3.9	0.5	not fire

```
In [31]: ### showing the unique values in the Classes feature  
df['Classes'].unique()
```

```
Out[31]: array(['not fire ', 'fire ', 'fire', 'fire ', 'not fire', 'not fire ',  
                'not fire ', nan, 'not fire '], dtype=object)
```

#### Observation:

- It seems that there are lot of spacing which should be stripped off.

```
In [32]: ### Stripping the object on both sides
```

```
for i in df.columns:  
    if df[i].dtypes=='O':  
        df[i]=df[i].str.strip()
```

```
In [33]: df['Classes'].unique()
```

```
Out[33]: array(['not fire', 'fire', nan], dtype=object)
```

#### Observation:

- There is some null values in the Classes feature which should be preprocessed

```
In [34]: ### There are one missing values which can be dropped off  
df['Classes'].isnull().sum()
```

```
Out[34]: 1
```



```
In [35]: df.isnull().sum()
```

```
Out[35]: day                0
month                0
Temperature          0
RH                  0
Ws                  0
Rain                0
FFMC                0
DMC                0
DC                  1
ISI                0
BUI                0
FWI                 1
Classes             1
dtype: int64
```

```
In [36]: df.dropna(inplace=True)
```

Observation: There are some missing values which can be dropped off from the dataframe

```
In [37]: df.isnull().sum()
```

```
Out[37]: day                0
month                0
Temperature          0
RH                  0
Ws                  0
Rain                0
FFMC                0
DMC                0
DC                  0
ISI                0
BUI                0
FWI                 0
Classes             0
dtype: int64
```

```
In [38]: df['Classes'].unique()
```

```
Out[38]: array(['not fire', 'fire'], dtype=object)
```

## What is Maximum Temperature When the Fire has happened

```
In [39]: df.groupby('Classes')['Temperature'].max()
```

```
Out[39]: Classes
fire      42
not fire   39
Name: Temperature, dtype: int32
```

Observation: The maximum temperature was around 42 degree celsius

***What was the Average Relative Humidity when the Fire happened where the Temperature was over 38 degree***

```
In [40]: df[(df['Classes']=='fire') & (df['Temperature']>38)][ 'RH'].mean()
```

```
Out[40]: 34.55555555555556
```

Observation: The average Humidity was around 34%

**Display all Average Temperature of month where the fire had happened**

```
In [41]: df[df['Classes']=='fire'].groupby(['month'])['Temperature'].mean()
```

```
Out[41]: month
6      32.440000
7      33.842105
8      35.647059
9      31.086957
Name: Temperature, dtype: float64
```

**Display all the details of the list where the fire has happened with maximum temperature**

```
In [42]: df[(df['Classes']=='fire') & (df['Temperature']==max(df['Temperature']))]
```

```
Out[42]:
```

	day	month	Temperature	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FWI	Classes
201	17	8	42	24	9	0.0	96.0	30.299999	76.400002	15.7	30.4	24.0	fire

**Display the Average Realtive Humidity for each month where temperature was more than 38 degree**

```
In [43]: df[df['Temperature']>38].groupby('month')['RH'].mean()
```

```
Out[43]: month
7      54.50
8      33.25
Name: RH, dtype: float64
```

**Display the skewness of the dataset**

```
In [44]: df.skew()
```

```
Out[44]: day      0.000365
month    -0.005207
Temperature -0.191327
RH       -0.242790
Ws        0.555586
Rain      4.568630
FFMC     -1.320130
DMC       1.522983
DC        1.473460
ISI       1.140243
BUI       1.452745
FWI       1.147593
dtype: float64
```

Observation: It seems that the rain is highly positive skewed has more outliers

### ***What is the Ram Space used by the dataset***

```
In [45]: df.memory_usage()
```

```
Out[45]: Index          1944
day             972
month           972
Temperature     972
RH              972
Ws              972
Rain            972
FFMC            972
DMC             972
DC              972
ISI             972
BUI             972
FWI             972
Classes         1944
dtype: int64
```

### ***What is the Temperature when Fire Weather Index was maximum***

```
In [46]: df[df['FWI']==max(df['FWI'])]
```

```
Out[46]:
```

	day	month	Temperature	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FWI	Classes
209	25	8	34	40	18	0.0	92.099998	56.299999	157.5	14.3	59.5	31.1	fire

### ***What is the Temperature when Fine Fuel Moisture Code was maximum***

```
In [47]: df[df['FFMC']==max(df['FFMC'])]
```

```
Out[47]:
```

	day	month	Temperature	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FWI	Classes
201	17	8	42	24	9	0.0	96.0	30.299999	76.400002	15.7	30.4	24.0	fire

### ***Covariance between Dependent and Independent Features***

In [48]: df.cov()

Out[48]:

	day	month	Temperature	RH	Ws	Rain	FFMC	DM
day	78.190729	-0.003639	3.119138	-9.969476	1.188603	-1.993174	28.544042	53.86313
month	-0.003639	1.242764	-0.229653	-0.681903	-0.124987	0.077762	0.272433	0.93867
Temperature	3.119138	-0.229653	13.162670	-35.043482	-2.901949	-2.372850	35.222858	21.83766
RH	-9.969476	-0.681903	-35.043482	219.874333	10.173809	6.604836	-137.215386	-75.07192
Ws	1.188603	-0.124987	-2.901949	10.173809	7.903887	0.965886	-6.718952	-0.02512
Rain	-1.993174	0.077762	-2.372850	6.604836	0.965886	4.012837	-15.634746	-7.16902
FFMC	28.544042	0.272433	35.222858	-137.215386	-6.718952	-15.634746	205.912202	107.34297
DMC	53.863133	0.938676	21.837668	-75.071928	-0.025120	-7.169025	107.342971	153.58743
DC	222.524339	6.722457	65.071727	-160.400449	10.604530	-28.456455	347.051365	517.42774
ISI	6.632060	0.303838	9.101371	-42.298446	0.099643	-2.891688	44.113113	35.03222
BUI	65.061368	1.349400	23.734918	-74.653741	1.257586	-8.546509	120.872509	173.20320
FWI	23.079143	0.685464	15.297068	-64.096917	0.677079	-4.835502	73.791798	80.76453

In [49]: *### Finding the outliers in the dataset*

```
def find_outliers_IQR(df):  
    q1=df.quantile(0.25)  
    q3=df.quantile(0.75)  
    IQR=q3-q1  
    lowerbound=q1-1.5*IQR  
    upperbound=q3+1.5*IQR  
    return lowerbound,upperbound
```

**Display all the Skewed data of Rain**

In [50]: lowerfence,upperfence=find\_outliers\_IQR(df['Rain'])

In [51]: print("The Lower Bound is",lowerfence,"\n"+"The Upper Bound is ",upperfence)

The Lower Bound is -0.75  
The Upper Bound is 1.25

In [52]: len(df[(df['Rain']<lowerfence)|(df['Rain']>upperfence)])

Out[52]: 35

Observation: There are 35 records in Rain which is skewed and has outliers which should be processed

**Display all the Skewed data of Duff Moisture Code**

In [53]: lowerfence,upperfence=find\_outliers\_IQR(df['DMC'])

```
In [54]: print("The Lower Bound is",lowerfence,"\n"+"The Upper Bound is ",upperfence)
```

```
The Lower Bound is -16.699999809265137
The Upper Bound is  43.30000019073486
```

```
In [55]: len(df[(df['DMC']<lowerfence)|(df['DMC']>upperfence)])
```


```
Out[55]: 12
```

Observation: There are 12 records in Duff Moisture Code which is skewed and has outliers which should be processed

## Graph Analysis

```
In [56]: graph_col=['Temperature','RH','Ws','Rain','FFMC','DMC','DC','ISI','BUI','FWI']
```

```
In [57]: ### Plotting the Kde Plot to check the distribution of the feature
fig, ax = plt.subplots(figsize=(5, 5))
for i in graph_col:
    sns.kdeplot(df[i])
    plt.xlabel(i)
    plt.ylabel('Distribution')
    plt.show()
```




### Observation:

1. Temperature is Normally Distributed.
2. Relative Humidity is slightly skewed right side.
3. Wind Speed is Right Skewed.
4. Rain is log normally distributed right skewed.
5. Fine Fuel Moisture Code is left skewed.
6. Duff Moisture Code (DMC), Drought Code (DC), Initial Spread Index (ISI), Build up Index (BUI), Fire Weather Index (FWI) are slightly right skewed.

**Display the Temperature when Fire has happened**

```
In [58]: fig, ax = plt.subplots(figsize=(10, 8))
sns.barplot(x='Classes',y='Temperature',data=df,palette="Blues_d")
```

```
Out[58]: <AxesSubplot:xlabel='Classes', ylabel='Temperature'>
```




**Observation:**

- The Average temperature was around 33 Celsius degrees when fire has happened

## Display the Distribution of Rain vs Temperature

```
In [59]: fig, ax = plt.subplots(figsize=(10, 8))
sns.boxplot(x="Temperature", y="Rain", data=df,
            whis=[0, 100], width=.6, palette="vlag")
```

Out[59]: <AxesSubplot:xlabel='Temperature', ylabel='Rain'>




**Observation:**

- During 28 Celsius degrees has maximum rainfall
- Temperature is inversely correlated with rain

## Show the Relation between Temperature and Fine Fuel Moisture Code With Respect to Classes

```
In [60]: fig, ax = plt.subplots(figsize=(10, 8))  
sns.scatterplot(x='Temperature',y='FFMC',hue='Classes',data=df)
```

```
Out[60]: <AxesSubplot:xlabel='Temperature', ylabel='FFMC'>
```



### Observation:


- Temperature and Fine Fuel Moisture Code are linearly increasing it is highly correlated
- The Temperature and FFMC are High when Fire has happened

## Show the relation between Temperature and Initial Spread Index



```
In [61]: fig, ax = plt.subplots(figsize=(10, 8))  
sns.scatterplot(x='Temperature', y='ISI', hue='Classes', data=df)
```

```
Out[61]: <AxesSubplot:xlabel='Temperature', ylabel='ISI'>
```




**Observation:**

- Temperature and Initial Spread Index are linearly increasing it is highly correlated
- The Temperature and ISI are High when Fire has happened

## MultiCollinearity

```
In [62]: fig, ax = plt.subplots(figsize=(10, 8))
sns.scatterplot(x='DMC',y='BUI',hue='Classes',data=df)
```

```
Out[62]: <AxesSubplot:xlabel='DMC', ylabel='BUI'>
```




**Observation:**

- Duff Moisture Code and Buildup Index are two independent variables which are highly correlated.
- Either one of the feature should be dropped for the feature selection.

**Show the relation between Temperature and Relative Humidity**

```
In [63]: fig, ax = plt.subplots(figsize=(10, 8))
sns.scatterplot(x='Temperature',y='RH',hue='Classes',data=df)
```

```
Out[63]: <AxesSubplot:xlabel='Temperature', ylabel='RH'>
```




**Observation:**

- Temperature and Relative Humidity are negative Correlated
- Rumidity is less when the fire has happened but the temperature was large

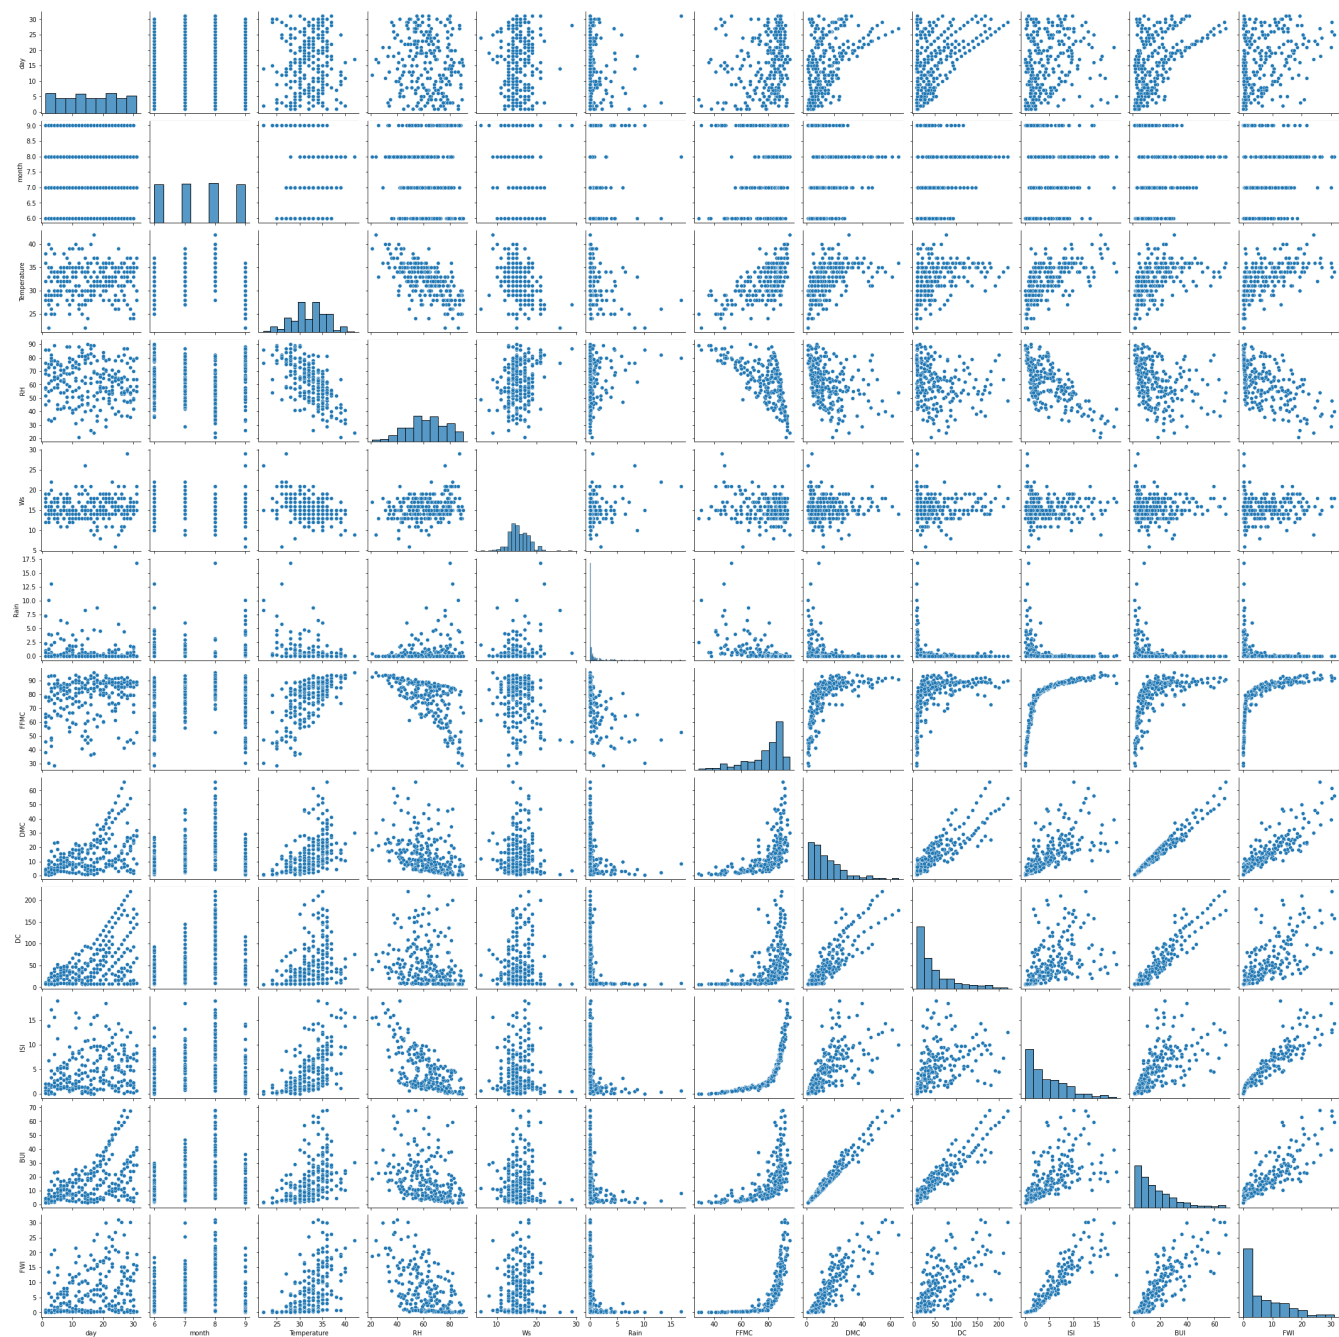
```
In [64]: fig, ax = plt.subplots(figsize=(10, 8))  
sns.barplot(x='month',y='Temperature',data=df,palette="Blues_d")
```

```
Out[64]: <AxesSubplot:xlabel='month', ylabel='Temperature'>
```



```
In [65]: sns.pairplot(df)
```


```
Out[65]: <seaborn.axisgrid.PairGrid at 0x21c04cab370>
```



## Display the Correlation Between the Features Gaphically

```
In [66]: fig, ax = plt.subplots(figsize=(12, 9))  
sns.heatmap(df.corr(),annot=True)
```

```
Out[66]: <AxesSubplot:>
```



### Observation:

- As Duff Moisture Code and Buildup Index are highly correlated there should be one feature since Duff Moisture Code is highly correlated with Temperature feature we can drop the Buildup Index

```
In [67]: df.drop(labels=['BUI'],axis=1,inplace=True)
```

```
In [68]: df.head()
```

```
Out[68]:
```

	day	month	Temperature	RH	Ws	Rain	FFMC	DMC	DC	ISI	FWI	Classes
0	1	6	29	57	18	0.0	65.699997	3.4	7.6	1.3	0.5	not fire
1	2	6	29	61	13	1.3	64.400002	4.1	7.6	1.0	0.4	not fire
2	3	6	26	82	22	13.1	47.099998	2.5	7.1	0.3	0.1	not fire
3	4	6	25	89	13	2.5	28.600000	1.3	6.9	0.0	0.0	not fire
4	5	6	27	77	16	0.0	64.800003	3.0	14.2	1.2	0.5	not fire

## Label Encoding

```
In [69]: from sklearn import preprocessing

# Label_encoder object knows how to understand word labels.
label_encoder = preprocessing.LabelEncoder()

# Encode labels in column 'Classes'.
df['Classes'] = label_encoder.fit_transform(df['Classes'])

df['Classes'].unique()
```

```
Out[69]: array([1, 0])
```

```
In [70]: df.head()
```

```
Out[70]:
```

	day	month	Temperature	RH	Ws	Rain	FFMC	DMC	DC	ISI	FWI	Classes
0	1	6	29	57	18	0.0	65.699997	3.4	7.6	1.3	0.5	1
1	2	6	29	61	13	1.3	64.400002	4.1	7.6	1.0	0.4	1
2	3	6	26	82	22	13.1	47.099998	2.5	7.1	0.3	0.1	1
3	4	6	25	89	13	2.5	28.600000	1.3	6.9	0.0	0.0	1
4	5	6	27	77	16	0.0	64.800003	3.0	14.2	1.2	0.5	1


### Observation:

-It changes all the Classes feature which has not fire to 1 and fire to 0

## How Linearly The Temperature and Fine Fuel Moisture Code are Related

```
In [71]: sns.set(rc={'figure.figsize':(12,9)})  
sns.regplot(x='FFMC',y='Temperature',data=df)
```

```
Out[71]: <AxesSubplot:xlabel='FFMC', ylabel='Temperature'>
```



**How Linearly The Temperature and Relative Humidity are Related**



```
In [72]: sns.set(rc={'figure.figsize':(12,9)})
sns.regplot(x='RH',y='Temperature',data=df)
```


```
Out[72]: <AxesSubplot:xlabel='RH', ylabel='Temperature'>
```



**How Linearly The Temperature and Initial Spread Index are Related**

```
In [73]: sns.set(rc={'figure.figsize':(15,10)})  
sns.regplot(x='ISI',y='Temperature',data=df)
```

```
Out[73]: <AxesSubplot:xlabel='ISI', ylabel='Temperature'>
```



```
In [74]: sample_temp=df['Temperature']
```

```
In [75]: df.drop(labels=['Temperature'],axis=1,inplace=True)
```

```
In [76]: df['Temperature']=sample_temp
```

```
In [77]: df.head()
```

```
Out[77]:
```

	day	month	RH	Ws	Rain	FFMC	DMC	DC	ISI	FWI	Classes	Temperature
0	1	6	57	18	0.0	65.699997	3.4	7.6	1.3	0.5	1	29
1	2	6	61	13	1.3	64.400002	4.1	7.6	1.0	0.4	1	29
2	3	6	82	22	13.1	47.099998	2.5	7.1	0.3	0.1	1	26
3	4	6	89	13	2.5	28.600000	1.3	6.9	0.0	0.0	1	25
4	5	6	77	16	0.0	64.800003	3.0	14.2	1.2	0.5	1	27

## Dependent and Independent Features

```
In [78]: X=df.iloc[:, :-1]  
y=df.iloc[:, -1]
```

```
In [79]: df['Classes'].unique()
```

```
Out[79]: array([1, 0])
```

```
In [80]: X.head()
```

```
Out[80]:
```

	day	month	RH	Ws	Rain	FFMC	DMC	DC	ISI	FWI	Classes
0	1	6	57	18	0.0	65.699997	3.4	7.6	1.3	0.5	1
1	2	6	61	13	1.3	64.400002	4.1	7.6	1.0	0.4	1
2	3	6	82	22	13.1	47.099998	2.5	7.1	0.3	0.1	1
3	4	6	89	13	2.5	28.600000	1.3	6.9	0.0	0.0	1
4	5	6	77	16	0.0	64.800003	3.0	14.2	1.2	0.5	1

```
In [81]: y.head()
```

```
Out[81]: 0    29  
1    29  
2    26  
3    25  
4    27  
Name: Temperature, dtype: int32
```

## Splitting the Dataset to Train and Test Feature

```
In [82]: from sklearn.model_selection import train_test_split
```

```
In [83]: X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.33,random_state=11)
```

```
In [84]: X_train.head()
```

```
Out[84]:
```

	day	month	RH	Ws	Rain	FFMC	DMC	DC	ISI	FWI	Classes
238	23	9	56	14	0.0	89.000000	29.400000	115.599998	7.5	15.2	0
233	18	9	33	13	0.1	90.599998	25.799999	77.800003	9.0	15.4	0
239	24	9	49	6	2.0	61.299999	11.900000	28.100000	0.6	0.4	1
198	14	8	40	13	0.0	91.900002	22.299999	55.500000	10.8	15.7	0
73	13	8	63	15	0.0	87.000000	19.000000	85.099998	5.9	10.2	0

```
In [85]: y_train.head()
```

```
Out[85]: 238    35
233    36
239    26
198    37
73     35
Name: Temperature, dtype: int32
```

```
In [86]: print("X_train shape is ",X_train.shape)
print("y_train shape is ",y_train.shape)
```

```
X_train shape is (162, 11)
y_train shape is (162,)
```

```
In [87]: print("X_test shape is ",X_test.shape)
print("y_test shape is ",y_test.shape)
```

```
X_test shape is (81, 11)
y_test shape is (81,)
```

## Standardizing the Independent Features

```
In [88]: from sklearn.preprocessing import StandardScaler
scaler=StandardScaler()
```

```
In [89]: X_train=scaler.fit_transform(X_train)
```

```
In [90]: X_test=scaler.transform(X_test)
```

## Model Training

```
In [91]: from sklearn.linear_model import LinearRegression
```

```
In [92]: regression=LinearRegression()
```

```
In [93]: regression
```

```
Out[93]: LinearRegression()
```

```
In [94]: regression.fit(X_train,y_train)
```

```
Out[94]: LinearRegression()
```

```
In [95]: ## Print the Coefficients and the intercept  
print(regression.coef_)
```

```
[-0.32525031 -0.55168719 -0.82572737 -0.44995213  0.25398265  1.28760797  
 0.25418955  0.98869337  0.73195835 -0.72209454 -0.05773034]
```

```
In [96]: print(regression.intercept_)
```

```
31.84567901234568
```


```
In [97]: ## Prediction for the test data  
regression_pred=regression.predict(X_test)
```

```
In [98]: regression_pred
```

```
Out[98]: array([30.26515065, 31.00882578, 33.97216377, 34.50657382, 30.68970807,  
 32.986536 , 34.35192609, 35.74002962, 33.98232054, 33.23717238,  
 31.83157616, 34.63147665, 32.21228429, 34.16809151, 34.15479705,  
 34.18021992, 30.56059095, 34.35811215, 32.11640042, 32.9586352 ,  
 32.60279639, 35.21682143, 26.30350593, 32.03320562, 28.0180797 ,  
 32.46962971, 31.17762763, 32.45781105, 28.25497111, 33.99194947,  
 31.70289223, 33.67761215, 33.58479649, 31.5898657 , 28.05247447,  
 31.96354957, 34.9390662 , 32.67067795, 36.2836806 , 27.22562961,  
 30.05900032, 25.67319687, 30.844162 , 28.59614825, 33.18342925,  
 25.24169735, 28.44700972, 27.87233434, 28.87002221, 25.92997604,  
 35.65904973, 35.59652034, 30.79326573, 35.94306315, 30.54203391,  
 33.10191871, 31.58229833, 33.35021373, 31.83475006, 31.39199406,  
 35.38699347, 31.77089059, 30.08144496, 34.73222826, 31.55068042,  
 34.07881214, 31.3204051 , 31.95762393, 34.91096036, 30.68004751,  
 29.09815236, 36.03760377, 32.68787322, 35.31710951, 33.22035314,  
 34.39110831, 32.11884417, 33.15635838, 35.98277233, 31.17914151,  
 33.12988661])
```

## Assumptions of Linear Regression

```
In [99]: plt.scatter(y_test, regression_pred)
plt.xlabel("Test Truth Data")
plt.ylabel("Test Predicted data")
plt.show()
```




```
In [100]: residuals=y_test-regression_pred
```

```
In [101]: residuals.head()
```

```
Out[101]: 39      2.734849  
         41     -0.008826  
         55      2.027836  
        165      1.493426  
        229     -2.689708  
         Name: Temperature, dtype: float64
```

```
In [102]: sns.displot(residuals,kind='kde')
```

```
Out[102]: <seaborn.axisgrid.FacetGrid at 0x21c0d4dcc10>
```



```
In [103]: plt.scatter(regression_pred,residuals)
```

```
Out[103]: <matplotlib.collections.PathCollection at 0x21c0d8a64f0>
```



## Calculating the Error

```
In [104]: from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_absolute_error
print("The Mean Squared Error for the model is",mean_squared_error(y_test,regression_pred))
print("The Mean Absolute Error for the model is",mean_absolute_error(y_test,regression_pred))
print("The Root Mean Squared Error for the model is",np.sqrt(mean_squared_error(y_test,regression_pred)))
```

The Mean Squared Error for the model is 6.651022710138245  
The Mean Absolute Error for the model is 2.0663555576296475  
The Root Mean Squared Error for the model is 2.5789576790126367

## Performance Metrics of the Model

```
In [105]: from sklearn.metrics import r2_score
score=r2_score(y_test,regression_pred)
print("The R2 Score for the model builded is",score)
```

The R2 Score for the model builded is 0.5014240665278434

```
In [106]: ## Adjusted R square
Adjusted_r=1-(1-score)*(len(y_test)-1)/(len(y_test)-X_test.shape[1]-1)
print("The Adjusted R Square for the model is",Adjusted_r)
```

The Adjusted R Square for the model is 0.4219409466989489

## Ridge Regression



```
In [107]: from sklearn.linear_model import Ridge  
ridge=Ridge()
```

```
In [108]: ridge.fit(X_train,y_train)
```

```
Out[108]: Ridge()
```

```
In [109]: ridge_pred=ridge.predict(X_test)
```


```
In [110]: ridge_pred
```

```
Out[110]: array([30.27356548, 30.98512337, 33.96056045, 34.50190218, 30.69818358,  
 32.99330427, 34.39672248, 35.76929768, 33.9903573 , 33.26028422,  
 31.80969456, 34.65959456, 32.20471906, 34.16037761, 34.16455094,  
 34.17485684, 30.56374068, 34.37662224, 32.11440336, 32.96285589,  
 32.58060596, 35.24585353, 26.34830107, 32.03104845, 28.05225902,  
 32.47166113, 31.17793816, 32.43674164, 28.28460747, 34.00311992,  
 31.65285399, 33.62550923, 33.60295552, 31.57666649, 28.08629548,  
 31.95054515, 34.940029 , 32.66164398, 36.35578088, 27.26689435,  
 30.02249231, 25.73115515, 30.81916108, 28.62802331, 33.2200882 ,  
 25.31952026, 28.40007621, 27.8933483 , 28.88994955, 25.98964768,  
 35.60150825, 35.65556054, 30.81227665, 36.03234329, 30.47859982,  
 33.09941955, 31.55413994, 33.35591305, 31.82429976, 31.38202554,  
 35.36968503, 31.75271506, 30.07992966, 34.74323182, 31.53419073,  
 34.06292577, 31.33569948, 31.95705101, 34.94254239, 30.68320829,  
 29.12267989, 36.07484378, 32.68418139, 35.316125 , 33.21917281,  
 34.39232191, 32.06728127, 33.12098996, 35.95235572, 31.14998099,  
 33.11007054])
```


```
In [111]: residual=ridge_pred-y_test
```

```
In [112]: sns.displot(residual,kind='kde')
```

```
Out[112]: <seaborn.axisgrid.FacetGrid at 0x21c0d8d4100>
```



```
In [113]: plt.scatter(y_test,ridge_pred)
plt.xlabel("Test Truth Data")
plt.ylabel("Test Predicted data")
plt.show()
```



## Calculating the Error

```
In [114]: from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_absolute_error
print("The Mean Squared Error for the model is",mean_squared_error(y_test,ridge_pred))
print("The Mean Absolute Error for the model is",mean_absolute_error(y_test,ridge_pred))
print("The Root Mean Squared Error for the model is",np.sqrt(mean_squared_error(y_test,ridge_pred)))
```

The Mean Squared Error for the model is 6.623472084700753  
The Mean Absolute Error for the model is 2.0602110174343062  
The Root Mean Squared Error for the model is 2.573610709625827

## Performance Metrics

```
In [115]: from sklearn.metrics import r2_score
score=r2_score(y_test,ridge_pred)
print("The R2 Score for the model builded is",score)
```

The R2 Score for the model builded is 0.5034893246684151

```
In [116]: ## Adjusted R square  
Adjusted_r=1-(1-score)*(len(y_test)-1)/(len(y_test)-X_test.shape[1]-1)  
print("The Adjusted R Square for the model is",Adjusted_r)
```

The Adjusted R Square for the model is 0.424335448890916

## Lasso Regression

```
In [117]: from sklearn.linear_model import Lasso  
lasso=Lasso()
```

```
In [118]: lasso.fit(X_train,y_train)
```

```
Out[118]: Lasso()
```

```
In [119]: lasso_pred=lasso.predict(X_test)
```


```
In [120]: lasso_pred
```

```
Out[120]: array([30.66221911, 31.22555026, 33.12098385, 33.33782614, 31.93070018,  
 32.75366151, 33.27869614, 33.82474232, 32.69335456, 32.99505173,  
 31.9295199 , 33.64251289, 32.10210686, 32.47762765, 33.10039985,  
 32.91693133, 30.88686771, 33.44393743, 31.95049949, 32.70849243,  
 32.31625281, 33.23971523, 27.02287838, 32.04826505, 29.95010242,  
 32.26116504, 31.64011077, 31.75577774, 29.32347015, 33.40699698,  
 30.97099365, 32.21029106, 32.84206539, 31.65859321, 30.18275653,  
 31.89979734, 33.12031663, 32.19947976, 33.88132322, 28.62951057,  
 30.88908738, 28.79966836, 31.57736497, 30.00753153, 32.53952522,  
 27.8658602 , 28.72621105, 28.96527201, 30.71402963, 27.74325415,  
 34.03609151, 33.65306912, 30.8793964 , 33.88904108, 30.7357015 ,  
 32.59182834, 31.77000172, 32.77293451, 31.87630076, 31.35379406,  
 32.93185331, 31.66122565, 29.65842251, 33.5361973 , 31.38261649,  
 33.14418373, 31.95752607, 32.02088091, 33.18338089, 31.8349038 ,  
 30.1304141 , 34.07974809, 32.42604777, 33.33360763, 32.37688954,  
 32.88082826, 30.73995372, 32.11262312, 34.26044048, 31.66800569,  
 32.28467754])
```

```
In [121]: residual=lasso_pred-y_test
```

```
In [122]: sns.displot(residual,kind='kde')
```

```
Out[122]: <seaborn.axisgrid.FacetGrid at 0x21c0f097e50>
```



## Calculation the Error

```
In [123]: from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_absolute_error
print("The Mean Squared Error for the model is",mean_squared_error(y_test,lasso_pred))
print("The Mean Absolute Error for the model is",mean_absolute_error(y_test,lasso_pred))
print("The Root Mean Squared Error for the model is",np.sqrt(mean_squared_error(y_test,
```

The Mean Squared Error for the model is 8.265102067478974  
The Mean Absolute Error for the model is 2.3820826872896768  
The Root Mean Squared Error for the model is 2.874909053775262

## Performance Metrics

```
In [124]: from sklearn.metrics import r2_score
score=r2_score(y_test,lasso_pred)
print("The R2 Score for the model builded is",score)
```

The R2 Score for the model builded is 0.38042897188508806

```
In [125]: ## Adjusted R square
Adjusted_r=1-(1-score)*(len(y_test)-1)/(len(y_test)-X_test.shape[1]-1)
print("The Adjusted R Square for the model is",Adjusted_r)
```

The Adjusted R Square for the model is 0.28165677899720354

## Elastic Net Regression

```
In [126]: from sklearn.linear_model import ElasticNet
```

```
In [127]: elastic = ElasticNet(random_state=0)
```

```
In [128]: elastic.fit(X_train,y_train)
```

```
Out[128]: ElasticNet(random_state=0)
```

```
In [129]: elastic_pred=elastic.predict(X_test)
```

```
In [130]: elastic_pred
```

```
Out[130]: array([30.32007329, 30.62849362, 33.44434918, 33.44559388, 31.60395251,
 32.76789206, 33.88217742, 34.5234838 , 33.21886391, 33.27182051,
 31.32321436, 34.2436066 , 31.88169277, 32.34655768, 33.44588338,
 32.87936362, 30.65260649, 33.75430321, 31.41332739, 32.68035549,
 32.24741161, 34.21345843, 27.63290233, 31.48134379, 29.715537 ,
 31.75659367, 31.0246927 , 31.18260669, 29.23899046, 33.69504347,
 30.47883983, 32.83177621, 32.79988727, 31.24890926, 29.80989008,
 31.61542413, 33.51619219, 31.75408187, 35.29538832, 28.63524758,
 30.60382831, 28.83784839, 30.68440236, 29.78478893, 33.15796974,
 28.12376044, 28.61876249, 28.91971221, 30.1411513 , 27.98833721,
 34.72538225, 34.8494712 , 30.5466458 , 35.4829652 , 30.36523242,
 32.63707799, 31.10475775, 32.66874449, 31.27404659, 31.11262502,
 33.94156321, 31.17736091, 29.55948724, 33.81517268, 31.04649904,
 33.45778954, 31.78279794, 31.73806541, 34.01686357, 31.3622155 ,
 29.83782929, 35.23320967, 32.04819372, 33.84195318, 32.81402498,
 32.9520279 , 30.5326578 , 32.34453875, 35.04514478, 30.89334041,
 32.44945411])
```


```
In [131]: residual=elastic_pred-y_test
```

```
In [132]: residual
```

```
Out[132]: 39    -2.679927
          41    -0.371506
          55    -2.555651
          65    -2.554406
          79    3.603953
          ...
          141   -2.467342
          59    1.344539
          196   -3.954855
          42    -3.106660
          57    -0.550546
          Name: Temperature, Length: 81, dtype: float64
```

```
In [133]: sns.displot(residual,kind='kde')
```

```
Out[133]: <seaborn.axisgrid.FacetGrid at 0x21c0f0ab400>
```



```
In [134]: from sklearn.metrics import mean_squared_error
          from sklearn.metrics import mean_absolute_error
          print("The Mean Squared Error for the model is",mean_squared_error(y_test,elastic_pred))
          print("The Mean Absolute Error for the model is",mean_absolute_error(y_test,elastic_pred))
          print("The Root Mean Squared Error for the model is",np.sqrt(mean_squared_error(y_test,elastic_pred)))
```

```
The Mean Squared Error for the model is 7.829970048889066
The Mean Absolute Error for the model is 2.2888159309658063
The Root Mean Squared Error for the model is 2.7982083640946156
```

```
In [135]: from sklearn.metrics import r2_score
          score=r2_score(y_test,elastic_pred)
          print("The R2 Score for the model builded is",score)
```

```
The R2 Score for the model builded is 0.4130474670860431
```

```
In [136]: ## Adjusted R square
          Adjusted_r=1-(1-score)*(len(y_test)-1)/(len(y_test)-X_test.shape[1]-1)
          print("The Adjusted R Square for the model is",Adjusted_r)
```

```
The Adjusted R Square for the model is 0.31947532415773117
```

