

Notebook changed



Individual household electric power consumption

The notebook file has changed on disk since the last time we opened or saved it. Do you want to overwrite the file on disk with the version open here, or load the version on disk (reload the page)?

Machine Learning Pipeline

- Data ingestion
- EDA
- Preprocessing
- Pickling for the preprocessing object(save the preprocessing model)

Cancel

Regression:linear regression,ridge regression,lasso regression,elastic net, support vector regression

Attribute Information:

Dataset Link:

<https://archive.ics.uci.edu/ml/datasets/Individual+household+electric+power+consumption>
(<https://archive.ics.uci.edu/ml/datasets/Individual+household+electric+power+consumption>)

- 1.date: Date in format dd/mm/yyyy
- 2.time: time in format hh:mm:ss
- 3.global_active_power: household global minute-averaged active power (in kilowatt)
- 4.global_reactive_power: household global minute-averaged reactive power (in kilowatt)
- 5.voltage: minute-averaged voltage (in volt)
- 6.global_intensity: household global minute-averaged current intensity (in ampere)
- 7.sub_metering_1: energy sub-metering No. 1 (in watt-hour of active energy). It corresponds to the kitchen, containing mainly a dishwasher, an oven and a microwave (hot plates are not electric but gas powered).
- 8.sub_metering_2: energy sub-metering No. 2 (in watt-hour of active energy). It corresponds to the laundry room, containing a washing-machine, a tumble-drier, a refrigerator and a light.
- 9.sub_metering_3: energy sub-metering No. 3 (in watt-hour of active energy). It corresponds to an electric water-heater and an air-conditioner.

```
In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import pymongo
from pymongo import MongoClient
%matplotlib inline
import warnings
warnings.filterwarnings("ignore")
```

```
In [3]: df=pd.read_csv(r'C:\Users\prasa\Desktop\Prasath\INeuron\INEURO~1\FSDS_B~1\FSDSBO~1\OCTM
```

```
In [4]: data=df.sample(n=50000,ignore_index=True)
```

In [5]: data.head()

Notebook changed

Out[5]:

	Date	Time	Global_active_power	Global_reactive_power	Voltage	Global_intensity	Sub_metering_1	Sub_metering_2	Sub_metering_3
0	9/3/2007	12:31:00	1.390	0.000	241.650	5.800	0.000	0.000	0.000
1	4/1/2009	01:51:00	0.482	0.096	247.020	2.000	0.000	0.000	0.000
2	12/11/2010	08:24:00	2.204	0.064	237.72	9.2	0.000	0.000	0.000
3	9/9/2007	22:26:00	1.030	0.368	239.880	4.600	0.000	0.000	0.000
4	13/7/2008	16:50:00	0.146	0.000	240.830	0.600	0.000	0.000	0.000

In [6]: data.shape

Out[6]: (50000, 9)

In [7]: data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50000 entries, 0 to 49999
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Date                   50000 non-null  object
1   Time                   50000 non-null  object
2   Global_active_power    50000 non-null  object
3   Global_reactive_power  50000 non-null  object
4   Voltage                50000 non-null  object
5   Global_intensity       50000 non-null  object
6   Sub_metering_1         50000 non-null  object
7   Sub_metering_2         50000 non-null  object
8   Sub_metering_3         49334 non-null  float64
dtypes: float64(1), object(8)
memory usage: 3.4+ MB
```

In [8]: data.isna().sum()

```
Out[8]: Date                0
Time                0
Global_active_power    0
Global_reactive_power  0
Voltage              0
Global_intensity      0
Sub_metering_1        0
Sub_metering_2        0
Sub_metering_3        666
dtype: int64
```

In [9]: data['Date'].unique()

```
Out[9]: array(['9/3/2007', '4/1/2009', '12/11/2010', ..., '8/1/2008', '13/8/2008',
               '12/10/2008'], dtype=object)
```

In [10]: data['Date']

Notebook changed

Out[10]: 0 9/3/2007

1 4/1/2009

2 12/11/2010

3 9/9/2007

4 13/7/2008

The notebook file has changed on disk since the last time we opened or saved it. Do you want to overwrite the file on disk with the version open here, or load the version on disk (reload the page)?

...

49995 12/3/2008

49996 21/7/2010

49997 8/9/2007

49998 7/1/2010

49999 29/5/2010

Name: Date, Length: 50000, dtype: object

Cancel

In [11]: data['Date']=pd.to_datetime(data['Date'])

In [12]: data['Year']=data['Date'].dt.year
data['Month']=data['Date'].dt.month

In [13]: data['Day']=data['Date'].dt.day

In [14]: data.head()

Out[14]:

	Date	Time	Global_active_power	Global_reactive_power	Voltage	Global_intensity	Sub_metering_1
0	2007-09-03	12:31:00	1.396	0.000	241.650	5.800	0.000
1	2009-04-01	01:51:00	0.482	0.096	247.020	2.000	0.000
2	2010-12-11	08:24:00	2.204	0.064	237.72	9.2	0.0
3	2007-09-09	22:26:00	1.030	0.368	239.880	4.600	0.000
4	2008-07-13	16:50:00	0.146	0.000	240.830	0.600	0.000



In [15]: data['Time'].unique()

Out[15]: array(['12:31:00', '01:51:00', '08:24:00', ..., '10:05:00', '08:22:00',
'21:49:00'], dtype=object)

In [16]: data['Time']=pd.to_datetime(data['Time'])

In [17]: data['Hour']=data['Time'].dt.hour

In [18]: data['Minutes']=data['Time'].dt.minute

In [19]: data.head()

Notebook changed

Out[19]:

	Date	Time	Global_active_power	Global_reactive_power	Voltage	Global_intensity	Sub_metering_1
The notebook file has changed on disk since the last time we opened or saved it. Do you want to overwrite the file on disk with the version opened here, or load the version on disk (reload the page)?							
0	2007-09-03	2022-11-04 12:31:00	1.396	0.000	241.650	5.800	0.000
1	2009-04-01	2022-11-04 01:51:00	0.482	0.096	247.020	2.000	0.000
2	2010-12-11	2022-11-04 08:24:00	2.204	0.064	237.72	9.2	0.0
3	2007-09-09	2022-11-04 22:26:00	1.030	0.368	239.880	4.600	0.000
4	2008-07-13	2022-11-04 16:50:00	0.146	0.000	240.830	0.600	0.000

In [20]: data.head()

Out[20]:

	Date	Time	Global_active_power	Global_reactive_power	Voltage	Global_intensity	Sub_metering_1
0	2007-09-03	2022-11-04 12:31:00	1.396	0.000	241.650	5.800	0.000
1	2009-04-01	2022-11-04 01:51:00	0.482	0.096	247.020	2.000	0.000
2	2010-12-11	2022-11-04 08:24:00	2.204	0.064	237.72	9.2	0.0
3	2007-09-09	2022-11-04 22:26:00	1.030	0.368	239.880	4.600	0.000
4	2008-07-13	2022-11-04 16:50:00	0.146	0.000	240.830	0.600	0.000

In [21]: data.info()

Notebook changed

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 50000 entries, 0 to 49999

Data columns (total 14 columns):

The notebook file has changed on disk since the last time we opened or saved it. Do you want to overwrite the file on disk with the version open here, or load the version on disk (reload the page)?

#	Column	Non-Null Count	Dtype
---	--------	----------------	-------

0	Date	50000 non-null	datetime64[ns]
1	Time	50000 non-null	datetime64[ns]
2	Global_active_power	50000 non-null	object
3	Global_reactive_power	50000 non-null	object
4	Voltage	50000 non-null	object
5	Global_intensity	50000 non-null	object
6	Sub_metering_1	50000 non-null	object
7	Sub_metering_2	50000 non-null	object
8	Sub_metering_3	49334 non-null	float64
9	Year	50000 non-null	int64
10	Month	50000 non-null	int64
11	Day	50000 non-null	int64
12	Hour	50000 non-null	int64
13	Minutes	50000 non-null	int64

dtypes: datetime64[ns](2), float64(1), int64(5), object(6)

memory usage: 5.3+ MB

In [22]: data_empty_Index=data[data['Global_active_power']=='?'].index

X

The notebook file has changed on disk since the last time we opened or saved it. Do you want to overwrite the file on disk with the version open here?

666 rows × 14 columns



```
data.replace(['?', 'nan', np.nan], -1, inplace=True)
```

In [25]: data.isna().sum()

Notebook changed

Out[25]:

Date	0
Time	0
Global_active_power	0
Global_reactive_power	0
Voltage	0
Global_intensity	0
Sub_metering_1	0
Sub_metering_2	0
Sub_metering_3	0
Year	0
Month	0
Day	0
Hour	0
Minutes	0

dtype: int64

The notebook file has changed on disk since the last time we opened or saved it. Do you want to overwrite the file on disk with the version open here, or load the version on disk (reload the page)?

Cancel

In [26]: data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50000 entries, 0 to 49999
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Date                  50000 non-null  datetime64[ns]
1   Time                  50000 non-null  datetime64[ns]
2   Global_active_power    50000 non-null  object
3   Global_reactive_power  50000 non-null  object
4   Voltage                50000 non-null  object
5   Global_intensity       50000 non-null  object
6   Sub_metering_1         50000 non-null  object
7   Sub_metering_2         50000 non-null  object
8   Sub_metering_3         50000 non-null  float64
9   Year                  50000 non-null  int64
10  Month                  50000 non-null  int64
11  Day                    50000 non-null  int64
12  Hour                   50000 non-null  int64
13  Minutes                50000 non-null  int64
dtypes: datetime64[ns](2), float64(1), int64(5), object(6)
memory usage: 5.3+ MB
```

In [27]: `from sklearn.impute import SimpleImputer`

Notebook changed

The notebook file has changed on disk since the last time we opened or saved it. Do you want to overwrite the file on disk with the version open here, or load the version on disk (reload the page)?

```
num_vars= ['Global_active_power', 'Global_reactive_power', 'Voltage',  
           'Global_intensity', 'Sub_metering_1', 'Sub_metering_2', 'Sub_metering_3']  
imp = SimpleImputer(missing_values=-1, strategy='mean')  
data[num_vars] = imp.fit_transform(data[num_vars])  
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 50000 entries, 0 to 49999  
Data columns (total 14 columns):  
#   Column                Non-Null Count  Dtype  
---  -  
0   Date                  50000 non-null  datetime64[ns]  
1   Time                  50000 non-null  datetime64[ns]  
2   Global_active_power    50000 non-null  float64  
3   Global_reactive_power  50000 non-null  float64  
4   Voltage                50000 non-null  float64  
5   Global_intensity       50000 non-null  float64  
6   Sub_metering_1         50000 non-null  float64  
7   Sub_metering_2         50000 non-null  float64  
8   Sub_metering_3         50000 non-null  float64  
9   Year                   50000 non-null  int64  
10  Month                  50000 non-null  int64  
11  Day                    50000 non-null  int64  
12  Hour                   50000 non-null  int64  
13  Minutes                50000 non-null  int64  
dtypes: datetime64[ns](2), float64(7), int64(5)  
memory usage: 5.3 MB
```

In [28]: `data.drop(['Date', 'Time'], axis=1, inplace=True)`

In [29]: `data.isna().sum()`

```
Out[29]: Global_active_power    0  
Global_reactive_power    0  
Voltage                  0  
Global_intensity         0  
Sub_metering_1           0  
Sub_metering_2           0  
Sub_metering_3           0  
Year                     0  
Month                    0  
Day                      0  
Hour                     0  
Minutes                  0  
dtype: int64
```

In [30]: `def merge_Metering(data):
 data['Total_Metering']=data['Sub_metering_1']+data['Sub_metering_2']+data['Sub_metering_3']
 return data`

In [31]: `data=merge_Metering(data)`

In [32]: data.head()

Notebook changed

Out[32]:

	Global_active_power	Global_reactive_power	Voltage	Global_intensity	Sub_metering_1	Sub_metering_2
0	1.396	0.000	241.65	5.8	0.0	0.0
1	0.482	0.096	247.02	2.0	0.0	0.0
2	2.204	0.064	237.72	9.2	0.0	0.0
3	1.030	0.368	239.88	4.6	0.0	1.0
4	0.146	0.000	240.83	0.6	0.0	0.0

In [33]: `# (global_active_power*1000/60 - sub_metering_1 - sub_metering_2 - sub_metering_3)`
`data['Power Consumed']=(data['Global_reactive_power']*1000)/60-data['Total_Metering']`

In [34]: data.head()

Out[34]:

	Global_active_power	Global_reactive_power	Voltage	Global_intensity	Sub_metering_1	Sub_metering_2
0	1.396	0.000	241.65	5.8	0.0	0.0
1	0.482	0.096	247.02	2.0	0.0	0.0
2	2.204	0.064	237.72	9.2	0.0	0.0
3	1.030	0.368	239.88	4.6	0.0	1.0
4	0.146	0.000	240.83	0.6	0.0	0.0

In [35]: data.describe().T

Out[35]:

	count	mean	std	min	25%	50%	75%	
Global_active_power	50000.0	1.093958	1.056850	0.076	0.310000	0.630	1.518000	9
Global_reactive_power	50000.0	0.123732	0.111672	0.000	0.050000	0.102	0.192000	
Voltage	50000.0	240.841642	3.207501	225.450	239.040000	240.970	242.830000	25
Global_intensity	50000.0	4.637885	4.442410	0.200	1.400000	2.800	6.400000	38
Sub_metering_1	50000.0	1.142721	6.138807	0.000	0.000000	0.000	0.000000	78
Sub_metering_2	50000.0	1.307577	5.858565	0.000	0.000000	0.000	1.000000	7
Sub_metering_3	50000.0	6.410670	8.356425	0.000	0.000000	1.000	17.000000	3
Year	50000.0	2008.436220	1.129122	2006.000	2007.000000	2008.000	2009.000000	2010
Month	50000.0	6.513720	3.436517	1.000	4.000000	7.000	9.000000	1
Day	50000.0	15.747520	8.848235	1.000	8.000000	16.000	23.000000	3
Hour	50000.0	11.546220	6.928231	0.000	6.000000	12.000	18.000000	2
Minutes	50000.0	29.510000	17.337448	0.000	14.000000	30.000	44.000000	5
Total_Metering	50000.0	8.860968	12.848745	0.000	0.000000	1.000	18.000000	12
Power Consumed	50000.0	-6.798762	12.649722	-120.800	-15.966667	0.000	1.333333	14

In [36]: data.corr()


Notebook changed

Out[36]:

	Global_active_power	Global_reactive_power	Voltage	Global_intensity	Sub_meterin
Global_active_power	1.000000	0.245048	-0.408372	0.998910	0.490378
Global_reactive_power	0.245048	1.000000	-0.110139	0.264056	0.129453
Voltage	-0.408372	-0.110139	1.000000	-0.419450	-0.204272
Global_intensity	0.998910	0.264056	-0.419450	1.000000	0.495463
Sub_metering_1	0.490378	0.129453	-0.204272	0.495463	1.000000
Sub_metering_2	0.439967	0.133839	-0.163492	0.444837	0.061000
Sub_metering_3	0.636064	0.085577	-0.274544	0.624219	0.106000
Year	-0.036677	0.040431	0.247092	-0.040535	-0.012000
Month	0.010138	0.014770	0.037184	0.009455	0.001000
Day	-0.014831	-0.002345	0.005556	-0.014965	-0.009000
Hour	0.279068	0.119727	-0.180074	0.279169	0.109000
Minutes	0.000104	-0.002247	0.014676	0.000036	-0.000000
Total_Metering	0.848576	0.178532	-0.350698	0.845522	0.575000
Power Consumed	-0.825872	-0.034206	0.340010	-0.819974	-0.565000

In [37]: sns.set(rc={'figure.figsize': (12,8)})
sns.histplot(data=data,x='Power Consumed',kde=True,bins=15)

Out[37]: <AxesSubplot:xlabel='Power Consumed', ylabel='Count'>



```
In [38]: sns.histplot(data=data,x='Global_active_power',kde=True,bins=15)
```

Notebook changed

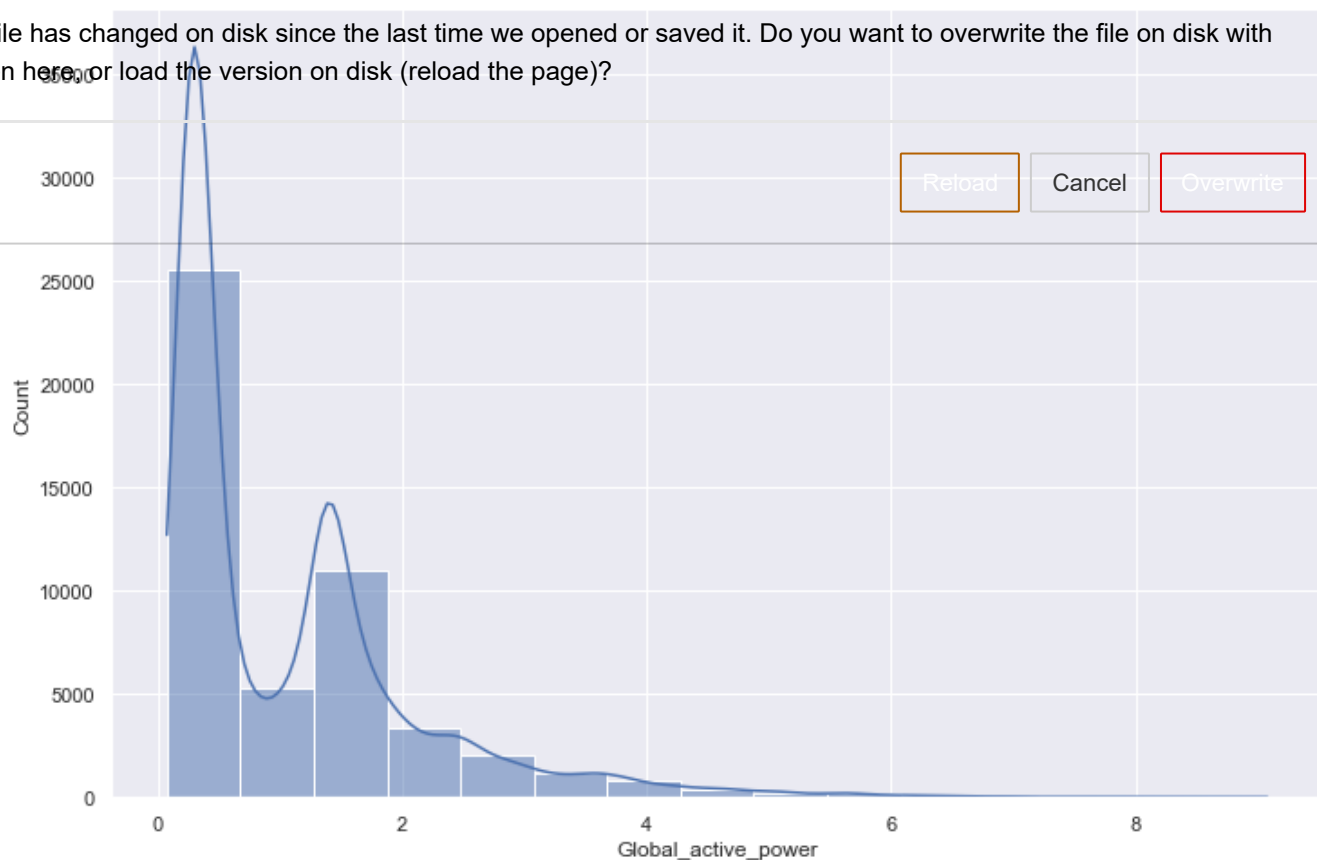
```
Out[38]: <AxesSubplot:xlabel='Global_active_power', ylabel='Count'>
```

The notebook file has changed on disk since the last time we opened or saved it. Do you want to overwrite the file on disk with the version open here, or load the version on disk (reload the page)?

Reload

Cancel

Overwrite



```
In [39]: sns.histplot(data=data,x='Total_Metering',kde=True,bins=15)
```

Notebook changed

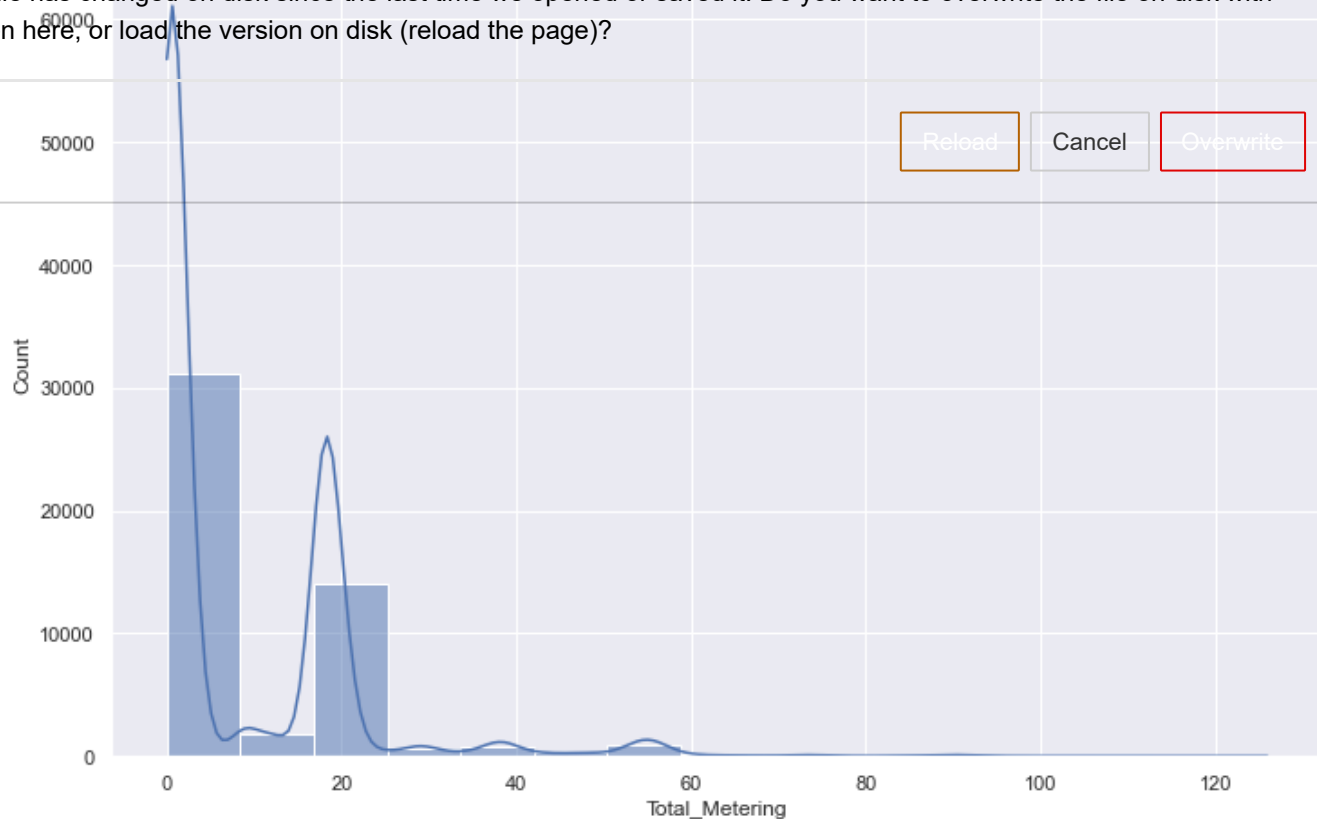
```
Out[39]: <AxesSubplot:xlabel='Total_Metering', ylabel='Count'>
```

The notebook file has changed on disk since the last time we opened or saved it. Do you want to overwrite the file on disk with the version open here, or load the version on disk (reload the page)?

Reload


Cancel

Overwrite



```
In [40]: sns.histplot(data=data,x='Global_intensity',kde=True,bins=15)
```

```
Out[40]: <AxesSubplot:xlabel='Global_intensity', ylabel='Count'>
```




```
In [41]: sns.histplot(data=data,x='Voltage',kde=True,bins=15)
```

Notebook changed

```
Out[41]: <AxesSubplot:xlabel='Voltage', ylabel='Count'>
```

The notebook file has changed on disk since the last time we opened or saved it. Do you want to overwrite the file on disk with the version open here, or load the version on disk (reload the page)?



Reload

Cancel


Overwrite

```
In [42]: sns.scatterplot(x='Power Consumed',y='Global_active_power',data=data,hue='Year')
```

Notebook changed


```
Out[42]: <AxesSubplot:xlabel='Power Consumed', ylabel='Global_active_power'>
```

The notebook file has changed on disk since the last time we opened or saved it. Do you want to overwrite the file on disk with the version open here, or load the version on disk (reload the page)?



```
In [43]: sns.scatterplot(x='Power Consumed',y='Total_Metering',data=data,hue='Year')
```


```
Out[43]: <AxesSubplot:xlabel='Power Consumed', ylabel='Total_Metering'>
```



In [44]: `sns.scatterplot(x='Power Consumed',y='Global_intensity',data=data,hue='Year')`
Notebook changed

Out[44]: `<AxesSubplot: xlabel='Power Consumed', ylabel='Global_intensity'>`

The notebook file has changed on disk since the last time we opened or saved it. Do you want to overwrite the file on disk with the version open here, or load the version on disk (reload the page)?




```
In [45]: sns.scatterplot(x='Power Consumed',y='Voltage',data=data,hue='Year')
```

Notebook changed


```
Out[45]: <AxesSubplot:xlabel='Power Consumed', ylabel='Voltage'>
```

The notebook file has changed on disk since the last time we opened or saved it. Do you want to overwrite the file on disk with the version open here, or load the version on disk (reload the page)?



```
In [46]: sns.scatterplot(x='Global_active_power',y='Global_intensity',data=data)
```

```
Out[46]: <AxesSubplot:xlabel='Global_active_power', ylabel='Global_intensity'>
```





```
In [47]: sns.scatterplot(x='Global_active_power',y='Total_Metering',data=data,hue='Year')
```

Notebook changed

```
Out[47]: <AxesSubplot:xlabel='Global_active_power', ylabel='Total_Metering'>
```

The notebook file has changed on disk since the last time we opened or saved it. Do you want to overwrite the file on disk with the version open here, or load the version on disk (reload the page)?



```
In [48]: data.head()
```

Out[48]:

	Global_active_power	Global_reactive_power	Voltage	Global_intensity	Sub_metering_1	Sub_metering_2
0	1.396	0.000	241.65	5.8	0.0	0.0
1	0.482	0.096	247.02	2.0	0.0	0.0
2	2.204	0.064	237.72	9.2	0.0	0.0
3	1.030	0.368	239.88	4.6	0.0	1.0
4	0.146	0.000	240.83	0.6	0.0	0.0

```
In [49]: data_1=data.drop(['Sub_metering_1','Sub_metering_2','Sub_metering_3','Year','Month','Day'])
```

```
In [50]: sns.heatmap(data_1.corr(),cmap="crest",annot=True)
Notebook changed
Out[50]: <AxesSubplot:>
```

The notebook file has changed on disk since the last time we opened or saved it. Do you want to overwrite the file on disk with the version open here, or load the version on disk (reload the page)?



```
In [51]: sns.violinplot(data=data_1 ,x="Global_active_power")
```

Notebook changed


```
Out[51]: <AxesSubplot:xlabel='Global_active_power'>
```

The notebook file has changed on disk since the last time we opened or saved it. Do you want to overwrite the file on disk with the version open here, or load the version on disk (reload the page)?

Reload

Cancel

Overwrite



```
In [52]: sns.violinplot(data=data_1 ,x="Global_intensity")
```

Notebook changed


```
Out[52]: <AxesSubplot:xlabel='Global_intensity'>
```

The notebook file has changed on disk since the last time we opened or saved it. Do you want to overwrite the file on disk with the version open here, or load the version on disk (reload the page)?

Reload

Cancel

Overwrite



```
In [53]: sns.violinplot(data=data_1,x='Total_Metering')
```

Notebook changed


```
Out[53]: <AxesSubplot:xlabel='Total_Metering'>
```

The notebook file has changed on disk since the last time we opened or saved it. Do you want to overwrite the file on disk with the version open here, or load the version on disk (reload the page)?

Reload

Cancel

Overwrite



```
In [54]: sns.violinplot(data=data_1,x='Voltage')
```

Notebook changed


```
Out[54]: <AxesSubplot:xlabel='Voltage'>
```

The notebook file has changed on disk since the last time we opened or saved it. Do you want to overwrite the file on disk with the version open here, or load the version on disk (reload the page)?

Reload

Cancel

Overwrite



```
In [55]: sns.regplot(x='Total_Metering',y='Power Consumed',data=data_1)
```

Notebook changed

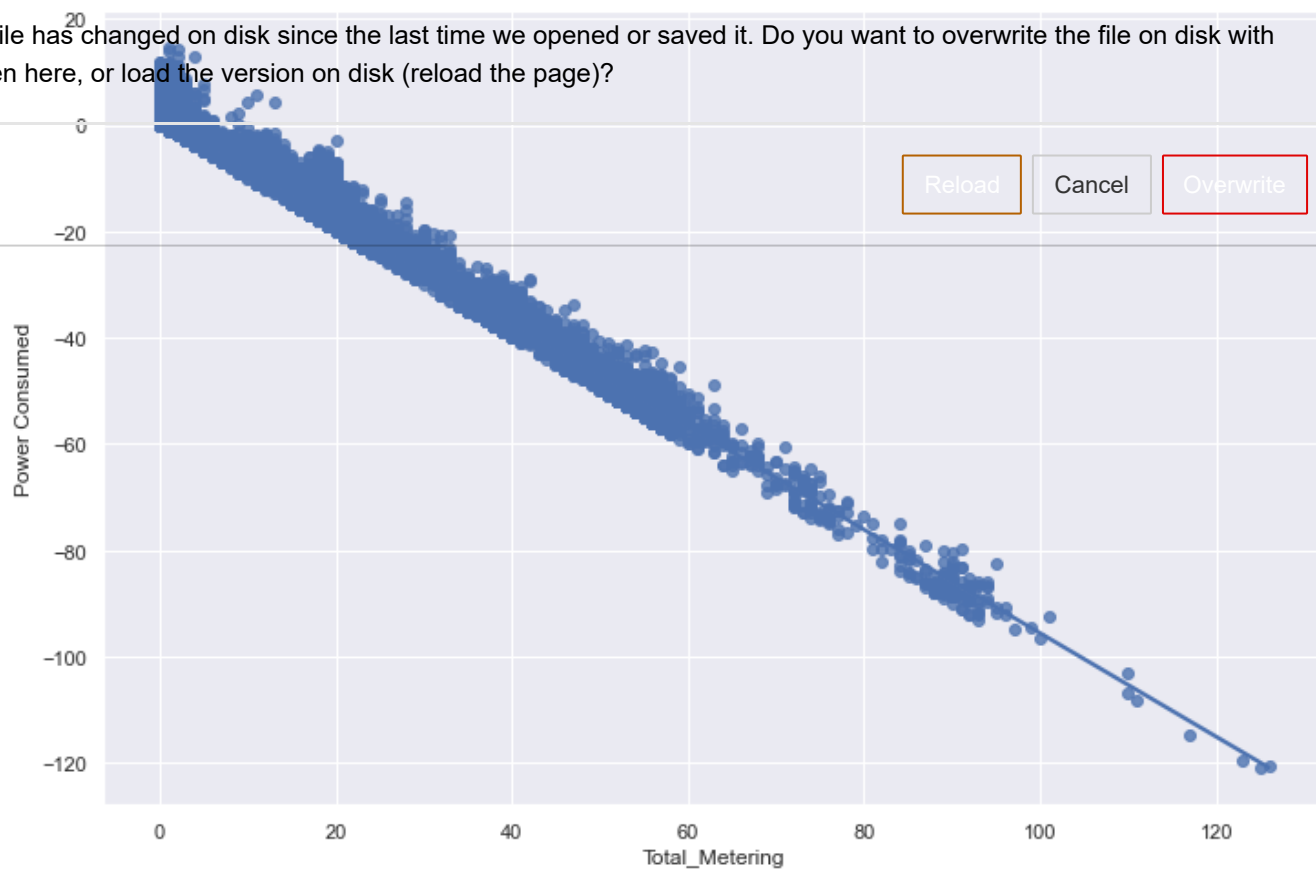
```
Out[55]: <AxesSubplot:xlabel='Total_Metering', ylabel='Power Consumed'>
```

The notebook file has changed on disk since the last time we opened or saved it. Do you want to overwrite the file on disk with the version open here, or load the version on disk (reload the page)?

Reload

Cancel

Overwrite

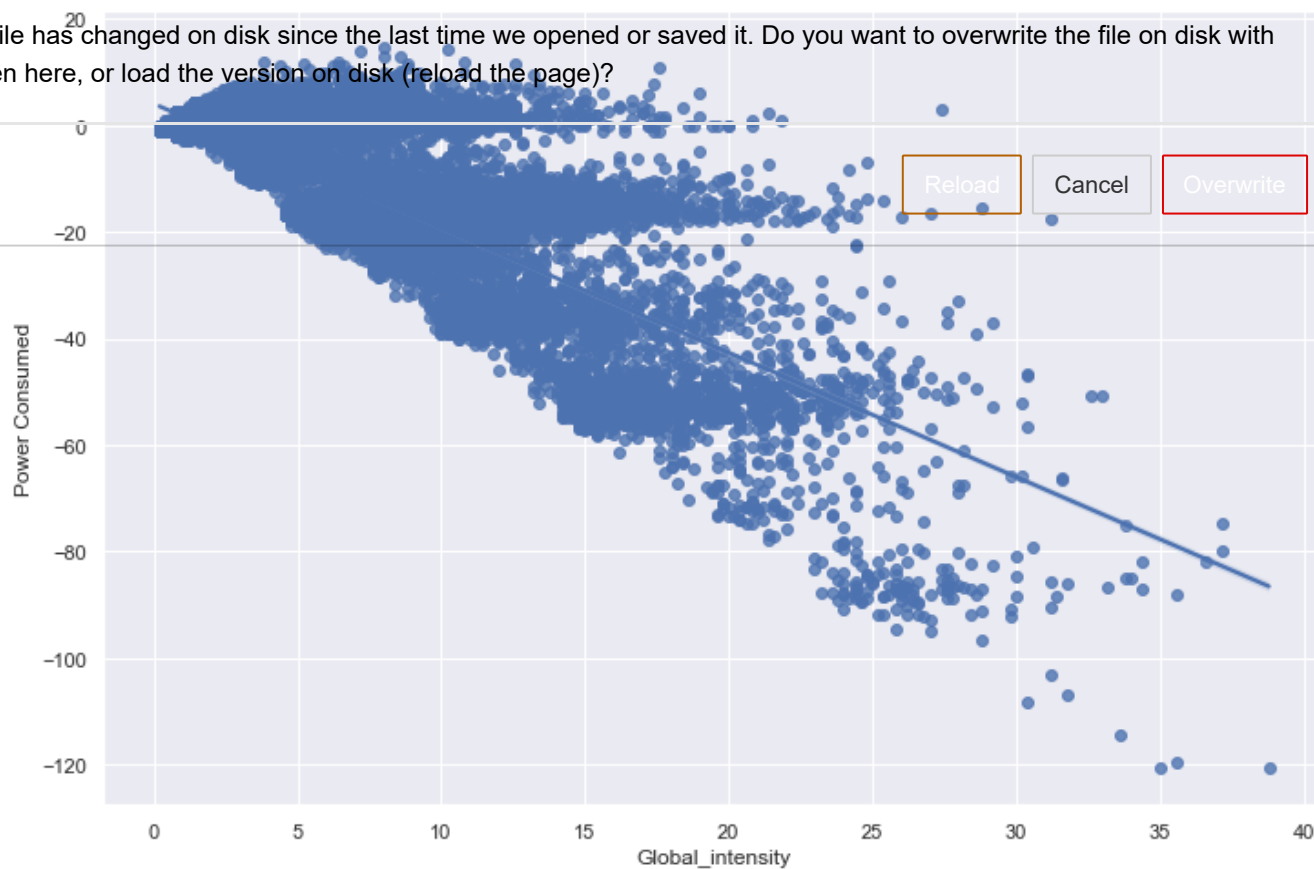


```
In [56]: sns.regplot(x='Global_intensity',y='Power Consumed',data=data_1)
```

Notebook changed

```
Out[56]: <AxesSubplot:xlabel='Global_intensity', ylabel='Power Consumed'>
```

The notebook file has changed on disk since the last time we opened or saved it. Do you want to overwrite the file on disk with the version open here, or load the version on disk (reload the page)?





```
In [57]: sns.regplot(x='Global_active_power',y='Power Consumed',data=data_1)
```

Notebook changed

```
Out[57]: <AxesSubplot:xlabel='Global_active_power', ylabel='Power Consumed'>
```

The notebook file has changed on disk since the last time we opened or saved it. Do you want to overwrite the file on disk with the version open here, or load the version on disk (reload the page)?

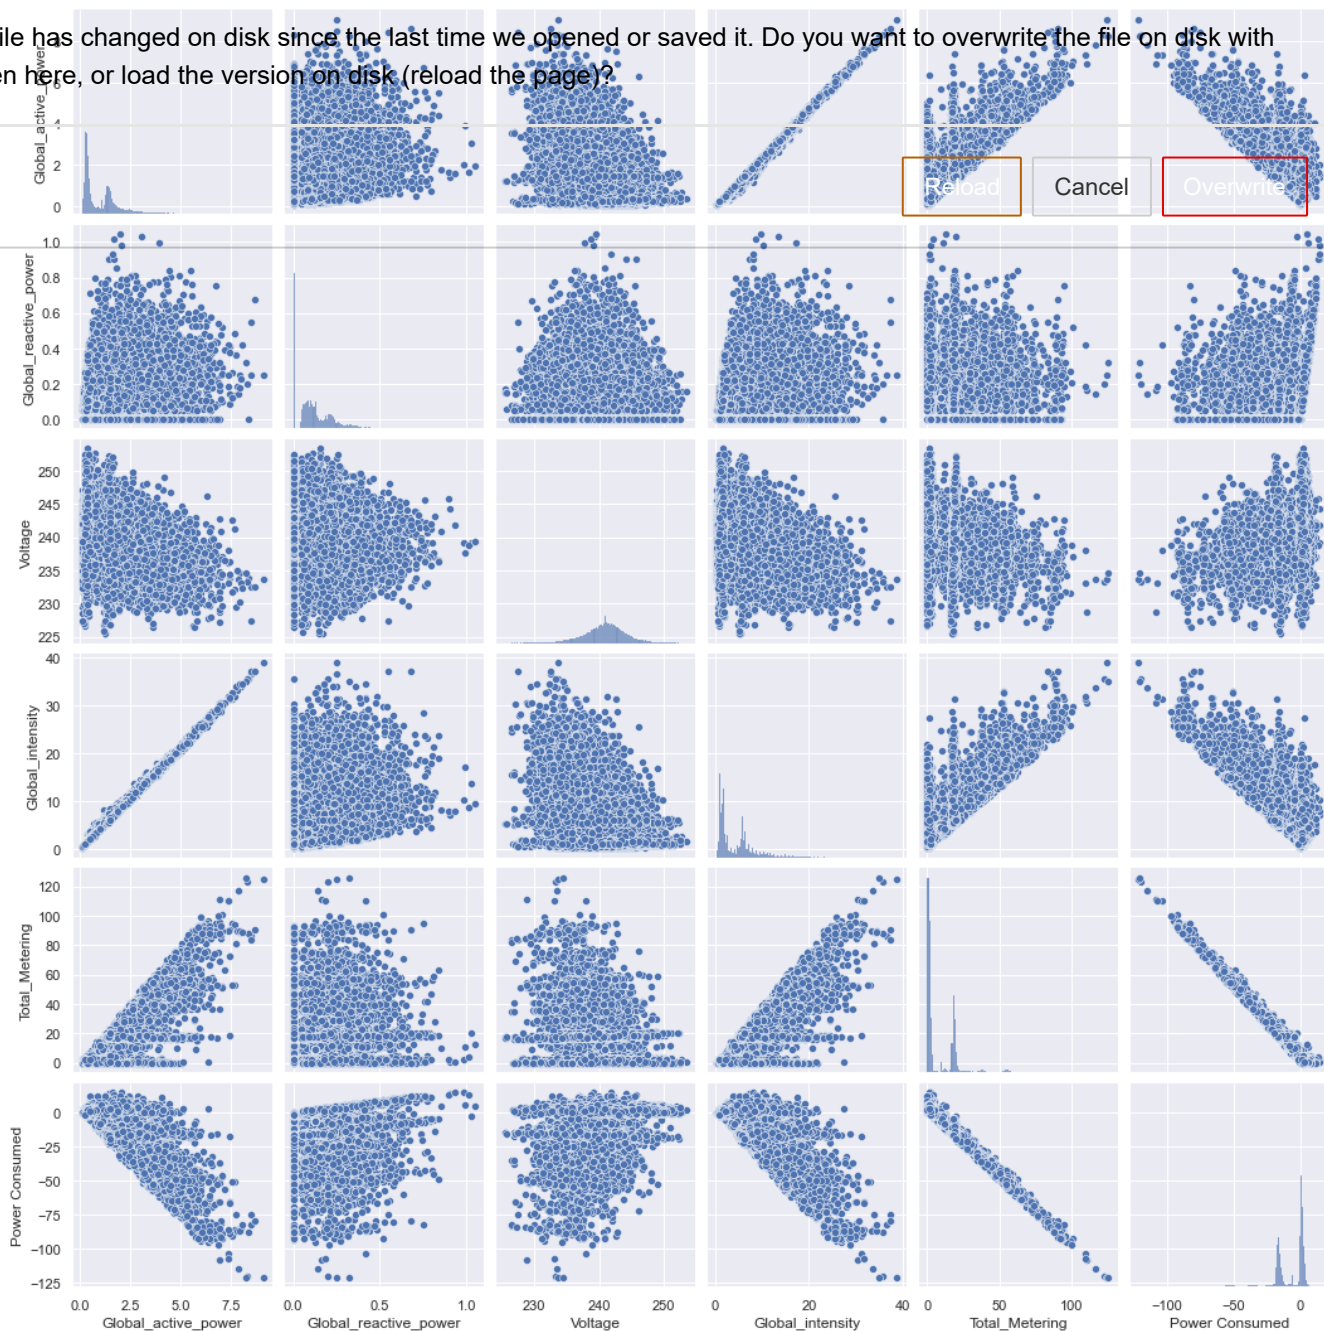


```
In [58]: sns.pairplot(data_1)
```

Notebook changed

```
Out[58]: <seaborn.axisgrid.PairGrid at 0x24c3be855b0>
```

The notebook file has changed on disk since the last time we opened or saved it. Do you want to overwrite the file on disk with the version open here, or load the version on disk (reload the page)?



```
In [106]: pd.to_pickle(data_1, 'Preprocessed.pkl')
```

In [59]: client = pymongo.MongoClient(
 Notebook changed mongodb+srv://prasathkps:Luci1108@cluster0.ldqcx.mongodb.net/?retryWrites=true&w=1
 db = client.test
 print(db)
 The notebook file has changed on disk since the last time we opened or saved it. Do you want to overwrite the file on disk with the version open here, or load the version on disk (reload the page)?
 Database(MongoClient(host=['cluster0-shard-00-02.ldqcx.mongodb.net:27017', 'cluster0-shard-00-01.ldqcx.mongodb.net:27017', 'cluster0-shard-00-00.ldqcx.mongodb.net:27017'], document_class=dict, tz_aware=False, connect=True, retrywrites=True, w='majority', aut
 hsource='admin', replicaset='atlas-2mebps-shard-0', tls=True), 'test') Cancel

In [60]: database=client['Power_Consumption_Data']
 collection=database['PowerConsumed']

In [107]: data1=pd.read_pickle('Preprocessed.pkl')

In [108]: df1=data_1.to_json('Power_Consumption.json')

In [109]: data_1.head()

Out[109]:

	Global_active_power	Global_reactive_power	Voltage	Global_intensity	Total_Metering	Power Consumed
0	1.396	0.000	241.65	5.8	18.0	-18.000000
1	0.482	0.096	247.02	2.0	0.0	1.600000
2	2.204	0.064	237.72	9.2	17.0	-15.933333
3	1.030	0.368	239.88	4.6	1.0	5.133333
4	0.146	0.000	240.83	0.6	0.0	0.000000

In [110]: df1=data_1.to_json('test1.json')

In [64]: df2=pd.read_json('test1.json')

In [65]: df2.head()

Out[65]:

	Global_active_power	Global_reactive_power	Voltage	Global_intensity	Total_Metering	Power Consumed
0	1.396	0.000	241.65	5.8	18.0	-18.000000
1	0.482	0.096	247.02	2.0	0.0	1.600000
2	2.204	0.064	237.72	9.2	17.0	-15.933333
3	1.030	0.368	239.88	4.6	1.0	5.133333
4	0.146	0.000	240.83	0.6	0.0	0.000000

In [66]: df3 = df2.to_dict(orient='records')

In [67]: print(type(df3))

<class 'list'>

In [68]: # collection.insert_many(df3)

```
In [69]: df = pd.DataFrame(list(collection.find({})))
```

Notebook changed

```
In [70]: df.head()
```

The notebook file has changed on disk since the last time we opened or saved it. Do you want to overwrite the file on disk with the version open here, or load the version on disk (reload the page)?

	_id	Global_active_power	Global_reactive_power	Voltage	Global_intensity	Total_Metering
0	63654da4b29a6c139caed5ba	1.078	0.320	245.43	4.6	1.0
1	63654da4b29a6c139caed5bb	1.786	0.250	240.64	7.4	1.0
2	63654da4b29a6c139caed5bc	1.926	0.260	236.61	8.2	0.0
3	63654da4b29a6c139caed5bd	1.820	0.122	236.06	7.6	19.0
4	63654da4b29a6c139caed5be	0.334	0.114	242.92	1.4	1.0

```
In [71]: df.drop('_id',axis=1,inplace=True)
```

```
In [72]: df.head()
```

Out[72]:

	Global_active_power	Global_reactive_power	Voltage	Global_intensity	Total_Metering	Power Consumed
0	1.078	0.320	245.43	4.6	1.0	4.333333
1	1.786	0.250	240.64	7.4	1.0	3.166667
2	1.926	0.260	236.61	8.2	0.0	4.333333
3	1.820	0.122	236.06	7.6	19.0	-16.966667
4	0.334	0.114	242.92	1.4	1.0	0.900000

```
In [73]: from sklearn.preprocessing import StandardScaler
scaler=StandardScaler()
scaler
```

Out[73]: StandardScaler()

```
In [74]: X=df.drop('Power Consumed',axis=1)
```

```
In [75]: y=df['Power Consumed']
```

```
In [76]: X.head()
```

Out[76]:

	Global_active_power	Global_reactive_power	Voltage	Global_intensity	Total_Metering
0	1.078	0.320	245.43	4.6	1.0
1	1.786	0.250	240.64	7.4	1.0
2	1.926	0.260	236.61	8.2	0.0
3	1.820	0.122	236.06	7.6	19.0
4	0.334	0.114	242.92	1.4	1.0

In [77]: y.head()

Notebook changed

Out[77]: 0 4.333333

1 3.166667

2 4.333333

3 -16.966667

4 0.900000

The notebook file has changed on disk since the last time we opened or saved it. Do you want to overwrite the file on disk with the version open here, or load the version on disk (reload the page)?

Name: Power Consumed, dtype: float64

Cancel

In [78]: from sklearn.model_selection import train_test_split

X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.33,random_state=42)

In [79]: print(X_train.shape,y_train.shape)

print(X_test.shape,y_test.shape)

(33500, 5) (33500,)

(16500, 5) (16500,)

In [80]: from sklearn.linear_model import LinearRegression

regression=LinearRegression()

regression

regression.fit(X_train,y_train)

Print the Coefficients and the intercept

print(regression.coef_)

print(regression.intercept_)

[-4.84729517e-13 1.66666667e+01 -3.81916720e-14 2.00728323e-13

-1.00000000e+00]

1.3847589741544652e-11

In [81]: *## Prediction for the test data*

regression_pred=regression.predict(X_test)

regression_pred

Out[81]: array([0.76666667, 0.9, -16.36666667, ..., 2.13333333,
1.6, 3.66666667])

In [82]: *# Calculating the Error*

from sklearn.metrics import mean_squared_error

from sklearn.metrics import mean_absolute_error

print("The Mean Squared Error for the model is",mean_squared_error(y_test,regression_pred))

print("The Mean Absolute Error for the model is",mean_absolute_error(y_test,regression_pred))

print("The Root Mean Squared Error for the model is",np.sqrt(mean_squared_error(y_test,regression_pred)))

The Mean Squared Error for the model is 2.6119828716863157e-21

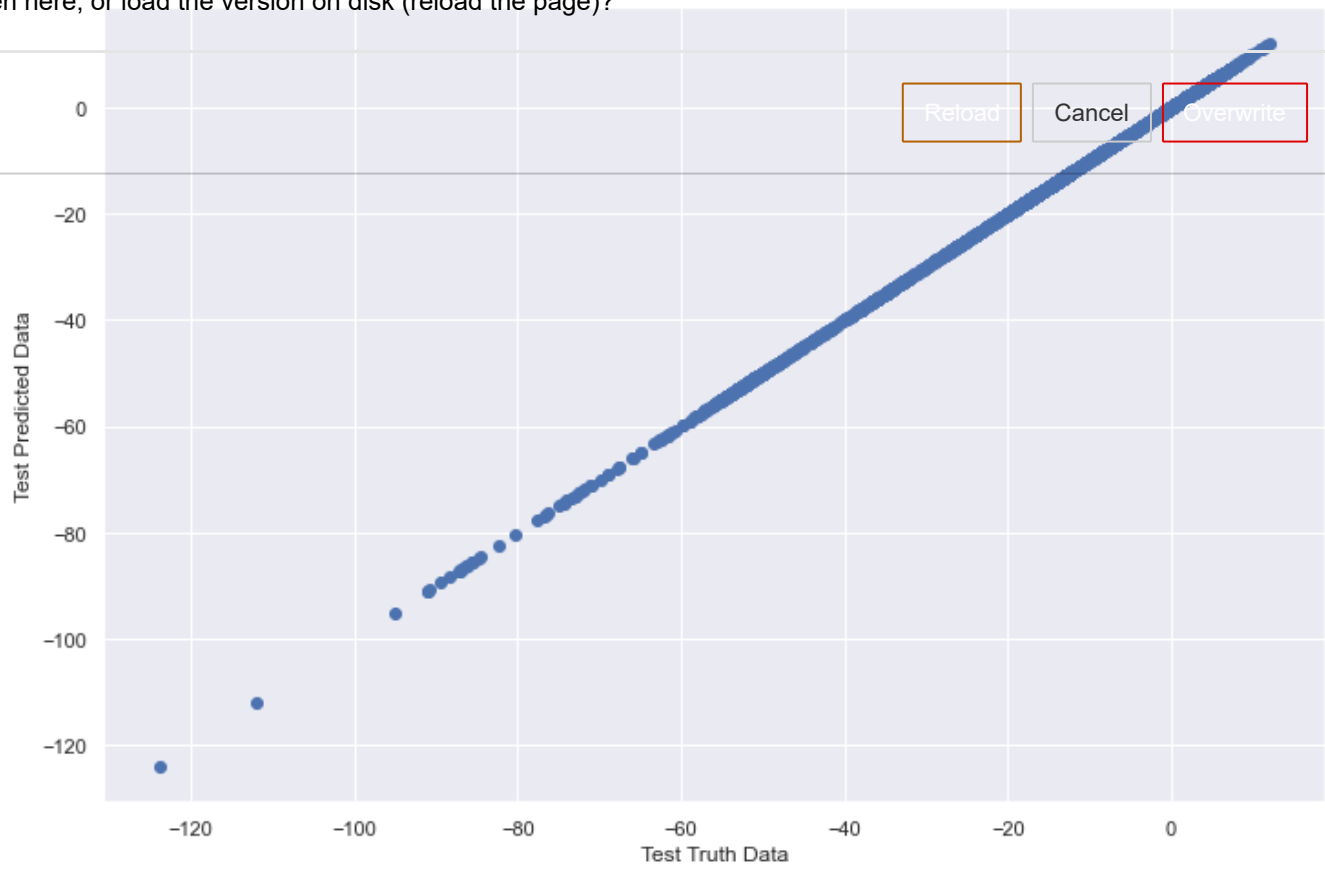
The Mean Absolute Error for the model is 2.4408202070745222e-11

The Root Mean Squared Error for the model is 5.1107561785770173e-11

```
In [83]: plt.scatter(y_test, regression_pred)
Notebook changed.
plt.xlabel("Test Truth Data")
plt.ylabel("Test Predicted Data")
```

```
Out[83]: Text(0, 0.5, 'Test Predicted Data')
```

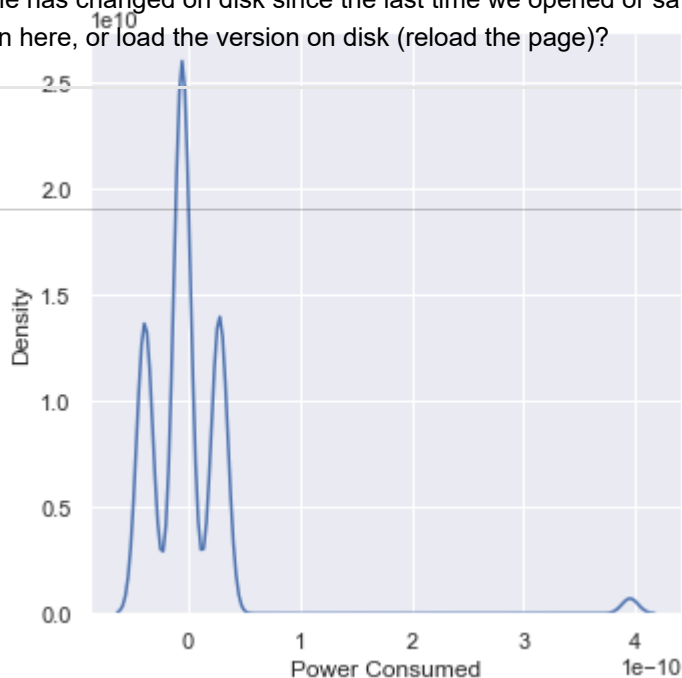
The notebook file has changed on disk since the last time we opened or saved it. Do you want to overwrite the file on disk with the version open here, or load the version on disk (reload the page)?



```
In [84]: residuals=y_test-regression_pred
Notebook changed. plot(residuals,kind="kde")
```

```
Out[84]: <seaborn.axisgrid.FacetGrid at 0x24c2d328370>
```

The notebook file has changed on disk since the last time we opened or saved it. Do you want to overwrite the file on disk with the version open here, or load the version on disk (reload the page)?



☐ ☐

```
In [85]: # Performance Metrics of the Model
from sklearn.metrics import r2_score
score=r2_score(y_test,regression_pred)
print("The R2 Score for the model builded is",score)
```

The R2 Score for the model builded is 1.0

```
In [86]: ## Adjusted R square
Adjusted_r=1-(1-score)*(len(y_test)-1)/(len(y_test)-X_test.shape[1]-1)
print("The Adjusted R Square for the model is",Adjusted_r)
```


The Adjusted R Square for the model is 1.0

```
In [87]: from sklearn.linear_model import Ridge
ridge=Ridge()
```

```
In [88]: ridge.fit(X_train,y_train)
Notebook changed. ridge_pred=ridge.predict(X_test)
```

The notebook file has changed on disk since the last time we opened or saved it. Do you want to overwrite the file on disk with the version open in the notebook (the version on disk) (reload the page)?

```
plt.scatter(y_test,ridge_pred)
plt.xlabel("Test Truth Data")
plt.ylabel("Test Predicted data")
plt.show()
```

```
In [90]: # Calculating the Error
from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_absolute_error
print("The Mean Squared Error for the model is",mean_squared_error(y_test,ridge_pred))
print("The Mean Absolute Error for the model is",mean_absolute_error(y_test,ridge_pred))
print("The Root Mean Squared Error for the model is",np.sqrt(mean_squared_error(y_test,
```

The Mean Squared Error for the model is 2.64423717500278e-05
The Mean Absolute Error for the model is 0.0038826204828712983
The Root Mean Squared Error for the model is 0.005142214673662293

```
In [91]: # Performance Metrics
from sklearn.metrics import r2_score
score=r2_score(y_test,ridge_pred)
print("The R2 Score for the model builded is",score)
```

The R2 Score for the model builded is 0.9999998271262797

```
In [92]: ## Adjusted R square
Adjusted_r=1-(1-score)*(len(y_test)-1)/(len(y_test)-X_test.shape[1]-1)
print("The Adjusted R Square for the model is",Adjusted_r)
```


The Adjusted R Square for the model is 0.9999998270738746

In [93]: `from sklearn.linear_model import Lasso`
Notebook changed

```
lasso = Lasso()  
lasso.fit(X_train, y_train)  
lasso_pred = lasso.predict(X_test)
```

The notebook file has changed on disk since the last time we opened or saved it. Do you want to overwrite the file on disk with the version open here, or load the version on disk (reload the page)?

Out[93]: <seaborn.axisgrid.FacetGrid at 0x24c2712d400>



In [94]: `# Calculation the Error`
`from sklearn.metrics import mean_squared_error`
`from sklearn.metrics import mean_absolute_error`
`print("The Mean Squared Error for the model is", mean_squared_error(y_test, lasso_pred))`
`print("The Mean Absolute Error for the model is", mean_absolute_error(y_test, lasso_pred))`
`print("The Root Mean Squared Error for the model is", np.sqrt(mean_squared_error(y_test, lasso_pred)))`

The Mean Squared Error for the model is 3.4289549294160917
The Mean Absolute Error for the model is 1.429987746272272
The Root Mean Squared Error for the model is 1.8517437537132646

In [95]: `# Performance Metrics`
`from sklearn.metrics import r2_score`
`score = r2_score(y_test, lasso_pred)`
`print("The R2 Score for the model builded is", score)`

The R2 Score for the model builded is 0.9775823363747858

In [96]: `## Adjusted R square`
`Adjusted_r = 1 - (1 - score) * (len(y_test) - 1) / (len(y_test) - X_test.shape[1] - 1)`
`print("The Adjusted R Square for the model is", Adjusted_r)`


The Adjusted R Square for the model is 0.9775755406722196

In [97]: `# Elastic Net Regression`
`from sklearn.linear_model import ElasticNet`
`elastic = ElasticNet(random_state=0)`
`elastic.fit(X_train,y_train)`

The notebook file has changed on disk since the last time it was opened or saved it. Do you want to overwrite the file on disk with the version open here, or load the version on disk (reload the page)?

`sns.displot(residual,kind='kde')`

Out[97]: <seaborn.axisgrid.FacetGrid at 0x24c2e7545b0>



In [98]: `from sklearn.metrics import mean_squared_error`
`from sklearn.metrics import mean_absolute_error`
`print("The Mean Squared Error for the model is",mean_squared_error(y_test,elastic_pred))`
`print("The Mean Absolute Error for the model is",mean_absolute_error(y_test,elastic_pred))`
`print("The Root Mean Squared Error for the model is",np.sqrt(mean_squared_error(y_test,elastic_pred)))`

The Mean Squared Error for the model is 3.381839955838527
The Mean Absolute Error for the model is 1.4224019688954508
The Root Mean Squared Error for the model is 1.838977965022563

In [99]: `from sklearn.metrics import r2_score`
`score=r2_score(y_test,elastic_pred)`
`print("The R2 Score for the model builded is",score)`
`## Adjusted R square`
`Adjusted_r=1-(1-score)*(len(y_test)-1)/(len(y_test)-X_test.shape[1]-1)`
`print("The Adjusted R Square for the model is",Adjusted_r)`

The R2 Score for the model builded is 0.9778903624792737
The Adjusted R Square for the model is 0.977883660151906

In [100]: `from sklearn.svm import SVR`
`svr=SVR()`
`svr`

Out[100]: SVR()

```
In [101]: svr.fit(X_train,y_train)
```

Notebook changed


```
Out[101]: SVR()
```

The notebook file has changed on disk since the last time we opened or saved it. Do you want to overwrite the file on disk with the version open here, or load the version on disk (reload the page)?

```
In [102]: svr_pred=svr.predict(X_test)
residual=svr_pred-y_test
sns.displot(residual,kind='kde')
```

```
Out[102]: <seaborn.axisgrid.FacetGrid at 0x24c4e059580>
```

Cancel



```
In [103]: from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_absolute_error
print("The Mean Squared Error for the model is",mean_squared_error(y_test,svr_pred))
print("The Mean Absolute Error for the model is",mean_absolute_error(y_test,svr_pred))
print("The Root Mean Squared Error for the model is",np.sqrt(mean_squared_error(y_test,svr_pred)))
```

The Mean Squared Error for the model is 3.515156979994131
The Mean Absolute Error for the model is 1.39604507906014
The Root Mean Squared Error for the model is 1.874875190511126

```
In [104]: from sklearn.metrics import r2_score
score=r2_score(y_test,svr_pred)
print("The R2 Score for the model builded is",score)
## Adjusted R square
Adjusted_r=1-(1-score)*(len(y_test)-1)/(len(y_test)-X_test.shape[1]-1)
print("The Adjusted R Square for the model is",Adjusted_r)
```

The R2 Score for the model builded is 0.977018768578346
The Adjusted R Square for the model is 0.9770118020355361