

Hyper Parameter Tuning GridSearchCV

- Data Ingestion
- EDA
- Feature Engineering
- Feature Scaling
- Model Training
- Performance Metrics
- Hyper parameter Tuning

```
In [1]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
data=pd.read_csv("https://raw.githubusercontent.com/aniruddhachoudhury/Red-Wine-Quality/
```

```
In [2]: data.head()
```

Out[2]:

| | fixed acidity | volatile acidity | citric acid | residual sugar | chlorides | free sulfur dioxide | total sulfur dioxide | density | pH | sulphates | alcohol | quality |
|---|------------------|---------------------|----------------|-------------------|-----------|---------------------------|----------------------------|---------|------|-----------|---------|---------|
| 0 | 7.4 | 0.70 | 0.00 | 1.9 | 0.076 | 11.0 | 34.0 | 0.9978 | 3.51 | 0.56 | 9.4 | 5 |
| 1 | 7.8 | 0.88 | 0.00 | 2.6 | 0.098 | 25.0 | 67.0 | 0.9968 | 3.20 | 0.68 | 9.8 | 5 |
| 2 | 7.8 | 0.76 | 0.04 | 2.3 | 0.092 | 15.0 | 54.0 | 0.9970 | 3.26 | 0.65 | 9.8 | 5 |
| 3 | 11.2 | 0.28 | 0.56 | 1.9 | 0.075 | 17.0 | 60.0 | 0.9980 | 3.16 | 0.58 | 9.8 | 6 |
| 4 | 7.4 | 0.70 | 0.00 | 1.9 | 0.076 | 11.0 | 34.0 | 0.9978 | 3.51 | 0.56 | 9.4 | 5 |

```
In [3]: data['quality'].value_counts()
```

```
Out[3]: 5    681
        6    638
        7    199
        4     53
        8     18
        3     10
        Name: quality, dtype: int64
```

```
In [4]: data['quality'].unique()
```

```
Out[4]: array([5, 6, 7, 4, 8, 3], dtype=int64)
```

```
In [5]: data.columns
```

```
Out[5]: Index(['fixed acidity', 'volatile acidity', 'citric acid', 'residual sugar',
               'chlorides', 'free sulfur dioxide', 'total sulfur dioxide', 'density',
               'pH', 'sulphates', 'alcohol', 'quality'],
              dtype='object')
```

```
In [6]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1599 entries, 0 to 1598
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fixed acidity          1599 non-null   float64
1   volatile acidity       1599 non-null   float64
2   citric acid            1599 non-null   float64
3   residual sugar         1599 non-null   float64
4   chlorides              1599 non-null   float64
5   free sulfur dioxide    1599 non-null   float64
6   total sulfur dioxide   1599 non-null   float64
7   density               1599 non-null   float64
8   pH                    1599 non-null   float64
9   sulphates              1599 non-null   float64
10  alcohol                1599 non-null   float64
11  quality                1599 non-null   int64
dtypes: float64(11), int64(1)
memory usage: 150.0 KB
```

```
In [7]: data.describe().T
```

Out[7]:

| | count | mean | std | min | 25% | 50% | 75% | max |
|-----------------------------|--------|-----------|-----------|---------|---------|----------|-----------|-----------|
| fixed acidity | 1599.0 | 8.319637 | 1.741096 | 4.60000 | 7.1000 | 7.90000 | 9.200000 | 15.90000 |
| volatile acidity | 1599.0 | 0.527821 | 0.179060 | 0.12000 | 0.3900 | 0.52000 | 0.640000 | 1.58000 |
| citric acid | 1599.0 | 0.270976 | 0.194801 | 0.00000 | 0.0900 | 0.26000 | 0.420000 | 1.00000 |
| residual sugar | 1599.0 | 2.538806 | 1.409928 | 0.90000 | 1.9000 | 2.20000 | 2.600000 | 15.50000 |
| chlorides | 1599.0 | 0.087467 | 0.047065 | 0.01200 | 0.0700 | 0.07900 | 0.090000 | 0.61100 |
| free sulfur dioxide | 1599.0 | 15.874922 | 10.460157 | 1.00000 | 7.0000 | 14.00000 | 21.000000 | 72.00000 |
| total sulfur dioxide | 1599.0 | 46.467792 | 32.895324 | 6.00000 | 22.0000 | 38.00000 | 62.000000 | 289.00000 |
| density | 1599.0 | 0.996747 | 0.001887 | 0.99007 | 0.9956 | 0.99675 | 0.997835 | 1.00369 |
| pH | 1599.0 | 3.311113 | 0.154386 | 2.74000 | 3.2100 | 3.31000 | 3.400000 | 4.01000 |
| sulphates | 1599.0 | 0.658149 | 0.169507 | 0.33000 | 0.5500 | 0.62000 | 0.730000 | 2.00000 |
| alcohol | 1599.0 | 10.422983 | 1.065668 | 8.40000 | 9.5000 | 10.20000 | 11.100000 | 14.90000 |
| quality | 1599.0 | 5.636023 | 0.807569 | 3.00000 | 5.0000 | 6.00000 | 6.000000 | 8.00000 |

In [8]: data.corr()

Out[8]:

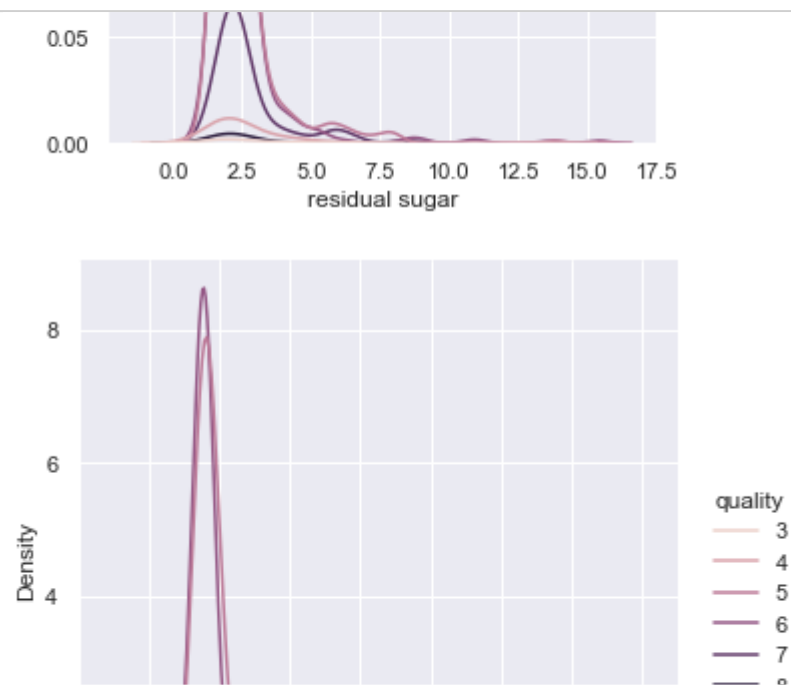
| | fixed acidity | volatile acidity | citric acid | residual sugar | chlorides | free sulfur dioxide | total sulfur dioxide | density | pH |
|----------------------|---------------|------------------|-------------|----------------|-----------|---------------------|----------------------|-----------|-----------|
| fixed acidity | 1.000000 | -0.256131 | 0.671703 | 0.114777 | 0.093705 | -0.153794 | -0.113181 | 0.668047 | -0.682978 |
| volatile acidity | -0.256131 | 1.000000 | -0.552496 | 0.001918 | 0.061298 | -0.010504 | 0.076470 | 0.022026 | 0.234937 |
| citric acid | 0.671703 | -0.552496 | 1.000000 | 0.143577 | 0.203823 | -0.060978 | 0.035533 | 0.364947 | -0.541904 |
| residual sugar | 0.114777 | 0.001918 | 0.143577 | 1.000000 | 0.055610 | 0.187049 | 0.203028 | 0.355283 | -0.085652 |
| chlorides | 0.093705 | 0.061298 | 0.203823 | 0.055610 | 1.000000 | 0.005562 | 0.047400 | 0.200632 | -0.265026 |
| free sulfur dioxide | -0.153794 | -0.010504 | -0.060978 | 0.187049 | 0.005562 | 1.000000 | 0.667666 | -0.021946 | 0.070377 |
| total sulfur dioxide | -0.113181 | 0.076470 | 0.035533 | 0.203028 | 0.047400 | 0.667666 | 1.000000 | 0.071269 | -0.066495 |
| density | 0.668047 | 0.022026 | 0.364947 | 0.355283 | 0.200632 | -0.021946 | 0.071269 | 1.000000 | -0.341699 |
| pH | -0.682978 | 0.234937 | -0.541904 | -0.085652 | -0.265026 | 0.070377 | -0.066495 | -0.341699 | 1.000000 |
| sulphates | 0.183006 | -0.260987 | 0.312770 | 0.005527 | 0.371260 | 0.051658 | 0.042947 | 0.148506 | -0.196648 |
| alcohol | -0.061668 | -0.202288 | 0.109903 | 0.042075 | -0.221141 | -0.069408 | -0.205654 | -0.496180 | 0.205633 |
| quality | 0.124052 | -0.390558 | 0.226373 | 0.013732 | -0.128907 | -0.050656 | -0.185100 | -0.174919 | -0.057731 |

In [9]: data.columns


Out[9]:

Index(['fixed acidity', 'volatile acidity', 'citric acid', 'residual sugar',
 'chlorides', 'free sulfur dioxide', 'total sulfur dioxide', 'density',
 'pH', 'sulphates', 'alcohol', 'quality'],
 dtype='object')


```
In [10]: sns.set(rc={"figure.figsize":(10, 7)})
for col in data.columns:
    sns.displot(data=data,x=data[col],hue='quality',kind='kde')
```



```
In [11]: sns.stripplot(data=data,y=data['alcohol'],x=data['quality'])
sns.set(rc={"figure.figsize":(10, 7)})
```




```
In [12]: sns.stripplot(data=data,y=data['total sulfur dioxide'],x=data['quality'])  
sns.set(rc={"figure.figsize":(10, 12)})
```




```
In [13]: sns.barplot(x='quality',data=data,y='alcohol')
```

```
Out[13]: <AxesSubplot:xlabel='quality', ylabel='alcohol'>
```



```
In [14]: sns.pairplot(data,hue="quality", diag_kind="hist")
```

```
Out[14]: <seaborn.axisgrid.PairGrid at 0x2367df8e0d0>
```



```
In [15]: data.head()
```

Out[15]:

| | fixed acidity | volatile acidity | citric acid | residual sugar | chlorides | free sulfur dioxide | total sulfur dioxide | density | pH | sulphates | alcohol | quality |
|---|---------------|------------------|-------------|----------------|-----------|---------------------|----------------------|---------|------|-----------|---------|---------|
| 0 | 7.4 | 0.70 | 0.00 | 1.9 | 0.076 | 11.0 | 34.0 | 0.9978 | 3.51 | 0.56 | 9.4 | 5 |
| 1 | 7.8 | 0.88 | 0.00 | 2.6 | 0.098 | 25.0 | 67.0 | 0.9968 | 3.20 | 0.68 | 9.8 | 5 |
| 2 | 7.8 | 0.76 | 0.04 | 2.3 | 0.092 | 15.0 | 54.0 | 0.9970 | 3.26 | 0.65 | 9.8 | 5 |
| 3 | 11.2 | 0.28 | 0.56 | 1.9 | 0.075 | 17.0 | 60.0 | 0.9980 | 3.16 | 0.58 | 9.8 | 6 |
| 4 | 7.4 | 0.70 | 0.00 | 1.9 | 0.076 | 11.0 | 34.0 | 0.9978 | 3.51 | 0.56 | 9.4 | 5 |

In [16]: `X=data.drop('quality',axis=1)`

In [17]: `X.head()`

Out[17]:

| | fixed acidity | volatile acidity | citric acid | residual sugar | chlorides | free sulfur dioxide | total sulfur dioxide | density | pH | sulphates | alcohol |
|---|---------------|------------------|-------------|----------------|-----------|---------------------|----------------------|---------|------|-----------|---------|
| 0 | 7.4 | 0.70 | 0.00 | 1.9 | 0.076 | 11.0 | 34.0 | 0.9978 | 3.51 | 0.56 | 9.4 |
| 1 | 7.8 | 0.88 | 0.00 | 2.6 | 0.098 | 25.0 | 67.0 | 0.9968 | 3.20 | 0.68 | 9.8 |
| 2 | 7.8 | 0.76 | 0.04 | 2.3 | 0.092 | 15.0 | 54.0 | 0.9970 | 3.26 | 0.65 | 9.8 |
| 3 | 11.2 | 0.28 | 0.56 | 1.9 | 0.075 | 17.0 | 60.0 | 0.9980 | 3.16 | 0.58 | 9.8 |
| 4 | 7.4 | 0.70 | 0.00 | 1.9 | 0.076 | 11.0 | 34.0 | 0.9978 | 3.51 | 0.56 | 9.4 |

In [18]: `y=data['quality']`

In [19]: `y.head()`

Out[19]:

| | |
|---|---|
| 0 | 5 |
| 1 | 5 |
| 2 | 5 |
| 3 | 6 |
| 4 | 5 |

Name: quality, dtype: int64

In [20]: `from sklearn.model_selection import train_test_split`
`X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.33,random_state=42)`

In [21]: `X_test.head()`

Out[21]:

| | fixed acidity | volatile acidity | citric acid | residual sugar | chlorides | free sulfur dioxide | total sulfur dioxide | density | pH | sulphates | alcohol |
|------|---------------|------------------|-------------|----------------|-----------|---------------------|----------------------|---------|------|-----------|---------|
| 803 | 7.7 | 0.56 | 0.08 | 2.50 | 0.114 | 14.0 | 46.0 | 0.9971 | 3.24 | 0.66 | 9.6 |
| 124 | 7.8 | 0.50 | 0.17 | 1.60 | 0.082 | 21.0 | 102.0 | 0.9960 | 3.39 | 0.48 | 9.5 |
| 350 | 10.7 | 0.67 | 0.22 | 2.70 | 0.107 | 17.0 | 34.0 | 1.0004 | 3.28 | 0.98 | 9.9 |
| 682 | 8.5 | 0.46 | 0.31 | 2.25 | 0.078 | 32.0 | 58.0 | 0.9980 | 3.33 | 0.54 | 9.8 |
| 1326 | 6.7 | 0.46 | 0.24 | 1.70 | 0.077 | 18.0 | 34.0 | 0.9948 | 3.39 | 0.60 | 10.6 |

In [22]: `X_train.head()`

Out[22]:

| | fixed acidity | volatile acidity | citric acid | residual sugar | chlorides | free sulfur dioxide | total sulfur dioxide | density | pH | sulphates | alcohol |
|------|---------------|------------------|-------------|----------------|-----------|---------------------|----------------------|---------|------|-----------|---------|
| 548 | 12.4 | 0.350 | 0.49 | 2.6 | 0.079 | 27.0 | 69.0 | 0.99940 | 3.12 | 0.75 | 10.4 |
| 355 | 6.7 | 0.750 | 0.01 | 2.4 | 0.078 | 17.0 | 32.0 | 0.99550 | 3.55 | 0.61 | 12.8 |
| 1296 | 6.6 | 0.630 | 0.00 | 4.3 | 0.093 | 51.0 | 77.5 | 0.99558 | 3.20 | 0.45 | 9.5 |
| 209 | 11.0 | 0.300 | 0.58 | 2.1 | 0.054 | 7.0 | 19.0 | 0.99800 | 3.31 | 0.88 | 10.5 |
| 140 | 8.4 | 0.745 | 0.11 | 1.9 | 0.090 | 16.0 | 63.0 | 0.99650 | 3.19 | 0.82 | 9.6 |


```
In [23]: y_train.head()
```

```
Out[23]: 548      6
          355      6
          1296     5
          209      7
          140      5
          Name: quality, dtype: int64
```

```
In [24]: y_test.head()
```

```
Out[24]: 803      6
          124      5
          350      6
          682      5
          1326     6
          Name: quality, dtype: int64
```

```
In [25]: from sklearn.preprocessing import StandardScaler
          scaler=StandardScaler()
          scaler
```

```
Out[25]: StandardScaler()
```

```
In [26]: scaler.fit(X_train)##calculate the mean and std dev
```

```
Out[26]: StandardScaler()
```

```
In [27]: print(scaler.mean_)
```

```
[ 8.30345472  0.53246499  0.26933707  2.54691877  0.08772736 15.91223156
 46.76330532  0.99677933  3.31453782  0.65881419 10.41521942]
```

```
In [28]: X_train_tf=scaler.fit_transform(X_train)
          X_train_tf
```

```
Out[28]: array([[ 2.40069523, -1.03103722,  1.12742595, ..., -1.26096312,
                  0.52726134, -0.01431863],
                [-0.93967131,  1.22920403, -1.32502245, ...,  1.52622836,
                 -0.28225704,  2.24363201],
                [-0.99827424,  0.55113165, -1.37611513, ..., -0.74241587,
                 -1.20742091, -0.86105011],
                ...,
                [-0.6466567 ,  0.49462562, -1.06955908, ...,  1.26695473,
                 -0.68701624, -0.86105011],
                [-0.23643625, -1.87862768,  0.4121285 , ...,  0.03540501,
                  0.81637505,  1.39690052],
                [-1.46709761, -1.3700734 , -0.04770558, ...,  0.48913386,
                 -0.68701624,  2.90220094]])
```

```
In [29]: X_test_tf=scaler.transform(X_test)
```

```
In [30]: from sklearn.svm import SVC
          model=SVC(kernel='linear',random_state=42)
          model
```

```
Out[30]: SVC(kernel='linear', random_state=42)
```

```
In [31]: model.fit(X_train_tf,y_train)
```

```
Out[31]: SVC(kernel='linear', random_state=42)
```

```
In [32]: model.score(X_train_tf,y_train)
```

```
Out[32]: 0.5994397759103641
```

```
In [33]: y_pred=model.predict(X_test_tf)
```

```
In [34]: y_test.head()
```

```
Out[34]: 803      6
         124      5
         350      6
         682      5
        1326      6
         Name: quality, dtype: int64
```

```
In [35]: from sklearn.metrics import accuracy_score
```

```
In [36]: accuracy_score(y_pred,y_test)
```

```
Out[36]: 0.5587121212121212
```

```
In [37]: from sklearn.model_selection import GridSearchCV
parameters=[{"C":[1,10,100,1000], 'kernel':['linear']}],
           {"C":[1,10,100,1000], 'kernel':['rbf'], 'gamma':[0.1,0.2,0.3,0.4,0.5,0.6,0.7,0.8]}]
gridSearch=GridSearchCV(estimator=model,param_grid=parameters,
                        scoring="accuracy",
                        cv=10,
                        n_jobs=-1)
gridSearch=gridSearch.fit(X_train_tf,y_train)
```

```
C:\Users\prasa\anaconda3\lib\site-packages\sklearn\model_selection\_split.py:676: User
Warning: The least populated class in y has only 8 members, which is less than n_split
s=10.
  warnings.warn(
```

```
In [38]: accuracy=gridSearch.best_score_
```

```
In [39]: accuracy
```

```
Out[39]: 0.6554690204222914
```

```
In [40]: gridSearch.best_params_
```

```
Out[40]: {'C': 10, 'gamma': 0.8, 'kernel': 'rbf'}
```

```
In [41]: model=SVC(C=10,kernel='rbf',gamma=0.8,random_state=42)
model
```

```
Out[41]: SVC(C=10, gamma=0.8, random_state=42)
```

```
In [42]: model.fit(X_train_tf,y_train)
```

```
Out[42]: SVC(C=10, gamma=0.8, random_state=42)
```

```
In [43]: y_pred=model.predict(X_test_tf)
```

```
In [44]: model.score(X_train_tf,y_train)
```

```
Out[44]: 0.9971988795518207
```

```
In [45]: from sklearn.metrics import accuracy_score  
accuracy=accuracy_score(y_test,y_pred)
```

```
In [46]: accuracy
```

```
Out[46]: 0.6193181818181818
```

```
In [47]: from sklearn.linear_model import LogisticRegression  
regressor=LogisticRegression()  
regressor
```

```
Out[47]: LogisticRegression()
```

```
In [48]: regressor.fit(X_train_tf,y_train)
```

```
Out[48]: LogisticRegression()
```

```
In [49]: y_pred2=regressor.predict(X_test_tf)
```

```
In [50]: regressor.score(X_train_tf,y_train)
```

```
Out[50]: 0.6218487394957983
```

```
In [51]: from sklearn.metrics import accuracy_score
```

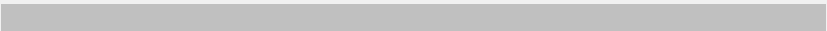
```
In [52]: accuracy_score(y_test,y_pred2)
```

```
Out[52]: 0.571969696969697
```

```
In [53]: import numpy as np  
C=np.linspace(1,100,10)  
penalty = ['l1', 'l2']  
C
```

```
Out[53]: array([ 1., 12., 23., 34., 45., 56., 67., 78., 89., 100.])
```

```
In [54]: solver_options = ['newton-cg', 'lbfgs', 'liblinear', 'sag']  
multi_class_options = ['ovr', 'multinomial']  
class_weight_options = ['None', 'balanced']  
param_grid = dict(C=C,penalty=penalty,solver = solver_options, multi_class = multi_class)
```

◀  ▶

```
In [55]: grid = GridSearchCV(estimator=regressor, param_grid=param_grid, cv=10, scoring = 'accuracy')
grid.fit(X_train_tf, y_train)
```

```
nan nan nan nan 0.51080824 nan
0.43702838 0.43702838 0.50987366 0.410964 nan nan
nan nan 0.40437002 0.40437002 nan 0.37541537
nan nan nan nan 0.60417965 0.60417965
nan 0.60417965 nan nan nan nan
nan nan nan nan nan nan
0.51174282 nan 0.43702838 0.43702838 0.50987366 0.38285739
nan nan nan nan 0.40437002 0.40437002
nan 0.38101419]
warnings.warn(
```

```
In [56]: grid.best_score_
```

```
Out[56]: 0.6041796469366563
```

```
In [57]: grid.best_params_
```

```
Out[57]: {'C': 1.0,
'class_weight': 'None',
'multi_class': 'ovr',
'penalty': 'l2',
'solver': 'newton-cg'}
```

```
In [58]: from sklearn.linear_model import LogisticRegression
regressor=LogisticRegression(C=5.5,class_weight='None',multi_class='ovr',solver='newton-cg')
regressor
```

```
Out[58]: LogisticRegression(C=5.5, class_weight='None', multi_class='ovr',
solver='newton-cg')
```

```
In [59]: regressor.fit(X_train_tf,y_train)
```

```
Out[59]: LogisticRegression(C=5.5, class_weight='None', multi_class='ovr',
solver='newton-cg')
```

```
In [60]: y_pred_Grid=regressor.predict(X_test_tf)
```

```
In [61]: accuracy_score(y_test,y_pred_Grid)
```

```
Out[61]: 0.5606060606060606
```